Software Engineering 350 Laboratory Project Manual: A Real-time Executive for Keil MCB1700

by

Yiqing Huang
Thomas Reidemeister

Electrical and Computer Engineering Department
University of Waterloo

Waterloo, Ontario, Canada, January 5, 2015

Contents

III Questions and Answers 73

List of Tables

List of Figures

Preface

The University of Waterloo Software Engineering (SE) SE350 course laboratory project is to design a small real-time executive (RTX) and implement it on a Keil MCB1700 board populated with an NXP LPC1768 microcontroller.

The main purpose of this document is a quick reference guide of the relevant hardware environment and software development tools of the Keil MCB1700 board for completing the laboratory project. To make the manual self-contained, we also include the project description [1] to further guide students.

There are two parts of the document.

- Part I: RTX Project Description

- Part III: Keil MCB1700 Reference Guide

Acknowledgments

We would like to sincerely thank Professor Paul Dasiewicz who originally designed the RTX course project and provided us with detailed notes and sample code. We also own many thanks to our students who did this course project in the past ten years and provided constructive feedback. Professor Sebastian Fischmeister made the Keil Boards and MDK-ARM donations possible. Professor Jim Barby provided timely departmental resource towards the development of the course project, without which this project will not be possible in winter 2012. Roger Sanderson oversees the ECE lab and provides us with all necessary experiment tools and resources, which we are grateful. We appreciate that Bernie Roehl has shared his valuable Keil board experiences with us. Our gratitude also goes out to Eric Praetzel who sets up the RTOS lab and also maintains the Keil software on Nexus machines; Laura Winger who managed to customize the boards so that we have the neat

---

[1]The original project description was written by Professor Paul Dasiewicz. The project description included in this manual is a modified version of the original one.

plastic cover to protect our hardware. Bob Boy from ARM always answers our questions in a detailed and timely manner. Thank everyone who has helped.

Lab Project Administration Policy

Project Group Policy

- **Group Size.** The project is done in groups of four. A group of less than four members is not recommended. There is no reduction in project deliverables regardless the size of the project group. Everyone in the group normally gets the same mark. The Course Book System at URL
  `https://ecewo32.uwaterloo.ca/cgi-bin/WebObjects/CourseBook`
  is used to signup for groups and reserve project demo times. *The project group signup is due by 4:30pm on the second Friday of the academic term.* Late group sign-up incurs a `5%` per day final lab mark deduction.

- **Group Split-up.** If you notice workload imbalance, try to solve it as soon as possible within your group or split-up the group as the last resort. Group split-up is only allowed once. There is one grace day deduction penalty to be applied to each member in the old group. We highly recommend everyone to stay with your group members as much as possible, for the ability to do team work will be an important skill in your future career. Please choose your lab partners carefully. A copy of the code and documentation completed before the group split-up will be given to each individual in the group.

- **Group Split-up Deadline.** To split from your group for a particular project deliverable, you need to notify the lab instructor in writing and sign the group slip-up form in the appendix at least one week before the particular project deliverable is due.

Project Submission Policy.

- **Project Submission and Due Dates.** The project is divided into four deliverables. Each deliverable requires the source code and a documentation file (in pdf file) [2]. Archive all files for each deliverable in a single file and submit it by using the Course Book System by performing group submission. Table 1 gives the weight, deadline and

---

[2]Put all source code (including all header files and binaries) and the documentation file in a separate directory. Include a README file with group identification, description of directory contents. Compress the directory contents into a single file. For archiving, you must choose zip.

Deliverable	Weight	Due Date	File Name
Group Sign-up		4:30pm Jan. 16th	
RTX P1	25%	4:30pm Feb. 3th	p1_Gid.zip
RTX P2	30%	4:30pm Mar. 10th	p2_Gid.zip
RTX P3	20%	4:30pm Mar. 24th	p3_Gid.zip
RTX P4	25%	4:30pm Apr. 6th	p4_Gid.zip

Table 1: Project Deliverable Weight and Deadlines. Replace the "id" in "Gid" with the two digit group ID

naming convention of each deliverable. A 15% penalty will be applied to a deliverable that is only able to function inside the simulator but not on the actual hardware.

- **Late Submissions.** There are *three grace days*[3] that can be used for project deliverables late submissions without incurring any penalty. A group split-up will consume one grace day. When you use up all your grace days, a 10% per day late penalty will be applied to a late submission. Please be advised that to simplify the book-keeping, late submission is counted in a unit of day rather than hour or minute. An hour late submission is one day late, so does a fifteen hour late submission. Unless notified otherwise, we always take the latest submission from the course book system. You are required to notify lab TAs (preferably by email) when you have used grace days.

Seeking Help Outside Scheduled Lab Hours

- **Discussion Forum.** We recommend students to use the Learn discussion forum to ask the teaching team questions instead of sending individual emails to lab teaching staff. For a question related to a deliverable, our target response time is one business day before the deadline of the particular deliverable [4]. *After the deadline, there is no guarantee on the response time.*

- **Office Hours.** During weeks where there are no scheduled labs, lab teaching staff hold bi-weekly office hours. Table 2 gives the office hour details.

- **Appointments.** Students can also make appointments with lab teaching staff should their problems are not resolved by discussion forum or during office hours. When

---

[3]Grace days are calendar days. Days in weekends are counted.

[4]Our past experiences show that the number of questions spike when deadline is close. The teaching staff will not be able to guarantee one business day response time when workload is above average, though we always try our best to provide timely response.

Time (even weeks only)	Location	Name	email ID
Wednesday 8:30-09:30	DC-2631	Yiqing (Irene) Huang	yqhuang
TBA			
TBA			

Table 2: Bi-weekly Office Hour Schedule

you request an appointment, please specify three preferred times and roughly how long you would like the appointment to be. On average, an appointment is fifteen minutes per project group.

Lab Facility After Hour Access Policy

After hour access to the lab will be given to the class when we start to use the Keil boards in lab. However please be advised that the after hour access is a privilege. Students are required to keep the lab equipment and furniture in good conditions to maintain this privilege.

No food or drink is allowed in the lab (water is permitted). Please be informed that you may share the lab with other classes. When resources become too tight, certain cooperation is required such as taking turns to use the stations in the lab.

Part I

RTX Project Description

Chapter 1

Introduction

1.1 Overview

In this project, you will design a small real-time executive (RTX) and implement it on a Keil MCB1700 board populated with an NXP LPC1768 microcontroller . The executive will provide a basic multiprogramming environment, with five priority levels, preemption, simple memory management, message-based inter-process communication, a basic timing service, system console I/O and debugging support.

Such an RTX is suitable for embedded computers which operate in real time. A cooperative, non-malicious software environment is assumed. The design of the RTX should allow its placement in ROM. Applications and non-kernel RTX processes must execute in the unprivileged level of LPC1768. The RTX kernel will execute in the privileged level [1]. It has 32K of RAM for use by the RTX and application processes. It contains four timers, four UARTs and several other peripheral interface devices. The board has two RS-232 interfaces, from which UART0 is used for your RTX system console and UART1 is used for your RTX debug terminal.

1.2 Summary of RTX Requirements

The summary of the RTX requirements are listed as follows:

---

[1]We do not require application processes to use Process Stack Pointer (PSP). You may use the Main Stack Pointer (MSP) both for your kernel and non-kernel code. However non-trivial implementations that are not required such as using PSP for user processes and using memory protection unit to safe guard kernel sensitive data will be rewarded with bonus marks.

1. **Scheduling Strategy**
   Four user priority levels plus an additional "hidden" priority level for the Null process, preemption, no time slicing, FIFO (First In, First Out) discipline at each priority level.

2. **RTX Primitives and Services**
   Refer to the Chapter 2 (Description of RTX Primitives and Services).

3. **RTX Footprint and Processor Loading**
   A reasonably *lean* implementation is expected. No standard C library function call is allowed in the kernel code.

4. **Error Detection and Recovery**
   At minimum, the RTX kernel must detect one type of error: an attempt to `send_message` to or `set_process_priority` of a non-existent process ID. The primitive will return an error code (a non-zero integer value). No error recovery is required. It may be assumed that the application processes can deal with this situation.

Chapter 2

Description of RTX Primitives and Services

This chapter lists the RTX primitive and services. You must implement theses as described and may not modify the prototypes in any way. The primitives listed below will always return a value, either a pointer or an `int` return code. In the latter case, the return code value of `0` indicates a success; non-zero value indicates a failure where applicable.

2.1 Memory Management

The RTX supports a simple memory management scheme. The memory is divided into blocks of fixed size (`128` bytes minimum). The size and the number of these blocks is a configuration parameter at compile time [1]. The blocks can be used by the requesting processes for storing local variables or as envelopes for messages sent to other processes. A block which is no longer needed must be returned to the RTX. Two primitives are to be provided.

```
void *request_memory_block();
```

The primitive returns a pointer to a memory block to the calling process. If no memory block is available, the calling process is blocked until a memory block becomes available. If several processes are waiting for a memory block and a block becomes available, the highest priority waiting process will get it.

---

[1] For example, you may define a macro as the number of memory blocks and change the macro value at compile time.

4

```
int release_memory_block(void *memory_block);
```

This primitive returns the memory block to the RTX. If there are processes waiting for a block, the block is given to the highest priority process, which is then unblocked. The caller of this primitive never blocks, but could be preempted. Thus, it may affect the currently executing process.

2.2 Processor Management

One primitive is to be provided.

```
int release_processor();
```

This primitive transfers the control to the RTX (the calling process voluntarily releases the processor). The invoking process remains ready to execute and is put at the end of the ready queue of the same priority. Another process may possibly be selected for execution.

2.3 Interprocess Communication

The RTX will support a message-based Interprocess Communication (IPC) discussed in lectures. Messages are carried in envelopes (memory blocks, see below) with a header which is less than 64 bytes. Two IPC primitives will be implemented.

```
int send_message(int process_id, void *message_envelope);
```

This primitive delivers to the destination process identified by `process_id` a message carried in a message envelope. The `message_envelope` argument is a pointer to caller-defined structure of the following general form:

```
struct msgbuf {
    int  mtype;   /* user defined message type */
    char mtext[1]; /* body of the message */
};
```

The `mtype` field takes a user defined message type. And the following macro defines the value of this field:

```
DEFAULT
```

A general purpose message.

```
KCD_REG
```

A message to register a command with the Keyboard Command Decoder Process (see Sectoin 3.1.2)

These macros are defined in `rtx.h` file as follows.

```
#define DEFAULT 0
#define KCD_REG 1
```

You are free to add more user defined message type macros.

The `mtext` field is an array (or other structure) whose size is limited to the size of one memory block less size of `type` field in the `msgbuf` structure. The primitive changes the state of destination process to ready-to-execute if appropriate. The sending process is preempted if the receiving process was blocked waiting for a message and has higher priority, otherwise the sender continues executing.

```
void *receive_message(int *sender_id);
```

This is a blocking receive. If there is a message waiting, a pointer to the message envelope containing it will be returned to the caller. If there is no such message, the calling process blocks and another process is selected for execution. The sender of the message is identified through `sender_id`, unless it is NULL. Note the `sender_id` is an output parameter and is not meant to filter which message to receive.

2.4   Timing Services

Unprivileged level processes obtain the timing service from RTX by the following primitive [2].

---

[2]Unprivileged processes should not read kernel timer data directly. You are free to add a primitive to return the kernel clock ticks to unprivileged processes should you find the delayed_send primitive is not sufficient to provide the timing service you need.

```
int delayed_send(int process_id, void *message_envelope, int delay);
```

The invoking process does not block. The message (in the memory block pointed to by the second parameter) will be sent to the destination process (`process_id`) after the expiration of the delay (timeout, given in *millisecond* units).

2.5   Process Priority

Process priorities have an integer priority value (0, 1, 2, 3, 4) where 0 is the highest priority level. Two primitives are to be provided to set and get the process priority.

```
int set_process_priority(int process_id, int priority);
```

This primitive sets the priority of the process with `process_id` to the value given in priority. A process may change priority of any process (including itself) except for i-processes (see Section 3.2). The priority of the null process may not be changed from level 4 and it is the only process that can be assigned to level 4 (see Section 3.1.1). The caller of this primitive never blocks, but could be preempted. This preemption may affect the currently executing process.

```
int get_process_priority(int process_id);
```

This primitive returns the priority of the process specified by the `process_id` parameter. For an invalid `process_id`, the primitive returns `-1`.

Chapter 3

Required Processes

This chapter describes the processes which you must implement for the project.

3.1 System Processes

System processes are those processes needed by the system to perform basic services (scheduling and I/O). You will need to make your design choice to determine which system processes require privileged level and which system processes may operate at unprivileged level.

3.1.1 The Null Process

This process runs as the lowest priority process (level 4) in the RTX. The Null process is the only process assigned to level 4. Level 4 is basically a"hidden" priority level reserved for the Null process. This preserves the four levels of user priorities (levels 0, 1, 2 and 3). Process id 0 is reserved for the null process. Initially, the following pseudo code can be used to design your null process:

```
loop forever
        release the processor
end loop
```

Once you have preemption working, then the "`release the processor`" line could be removed from the infinite loop.

3.1.2   System Console I/O Processes

The system console is used for communication with the RTX and application processes. It consists of two devices: keyboard and CRT display. These two devices communicate serially with the microcomputer; using the receive and transmit lines of one of the two RS-232 ports.

The RTX will include two system processes, the Keyboard Command Decoder (KCD) process and the CRT Display process. These processes work in cooperation with the UART interrupt handler i-process.

The Keyboard Command Decoder (KCD) Process

A keyboard command starts with the prompt character %, followed by a single (or multiple) letter command identifier and possibly additional command data. For example, `%WS 12:45:00` could be a command to the wall clock process, telling it to start the wall clock display and setttig the current time to `12:45:00` (where the command format is `%WS hh:mm:ss`). The command decoder process responds to two types of messages: console keyboard input and command registra-tion. The latter contains the command identifier and the process id of the process to which such commands are to be delivered when entered on the console keyboard. The processing of messages received depends on their type:

- Command Registration
  The command identifier is associated with the process id of the registrant.

- Keyboard Input
  The string input is sent to CRT display for output. If the string begins with a registered command identifier, it is also sent to the registered requester.

The CRT Display Process

This process responds to only one message type: a CRT display request. The message body contains the character string to be displayed. The string may contain control characters (e.g. newline). The process causes the string to be output to the console CRT. In printing to the console display, the process must use the UART i-process. Any message received is freed using the `release_memory_block` primitive.

3.2   Interrupt Processes (I-Processes)

Two interrupt handling processes are required:

3.2.1 The Timer I-Process

The timer i-process is executed each time a hardware timer interrupt occurs. The timer i-process should handle the delivery of delayed send messages after the required time has expired.

3.2.2 The UART I-Process

The UART i-process uses interrupts for both the transmission and receiving of characters from the serial port. No polling or busy waiting strategies may be implemented. The UART i-process forwards characters (or commands) received to the KCD, and also responds to messages received from the CRT display process to transmit characters to the serial port.

Three Hot Keys

As well, the UART i-process is used to provide debugging services which will be used during the demonstration. Upon receiving a specific character (a hot key - your choice, e.g., !) as input, the UART i-process will print one of the following to the RTX system debug terminal.

- The processes currently on the ready queue(s) and their priority.

- The processes currently on the blocked on memory queue(s) and their priorities.

- The processes currently on the blocked on receive queue(s) and their priorities.

As well, you are free to implement other hot keys to help in debugging. For example, a hot key which lists the processes, their priorities, their states; or another which prints out the number of memory blocks available. Like all other debug prints, the hot key implementation should be wrapped in

```
#ifdef _DEBUG_HOTKEYS
...
#endif
```

preprocessor statements and should be turned off during automated testing. If the automated test processes fail, you may be asked to turn the hot keys on again in determining why the test processes are failing. Another hot key debug printout may be used to display recent interprocess message passing. A (circular) log buffer keeps track of the 10 most

recent `send_message` and `receive_message` invocations made by the processes; upon receiving a specific hot key, these most recent 10 sent and 10 received messages are printed to the debug con-sole. The number 10 is used only as an example. The information printed could contain information such as:

- Sender process id

- Destination process id

- Message type

- First 16 bytes of the message text

- The time stamp of the transaction (using the RTX clock)

3.3   User Processes

These processes operate at *unprivileged level* and will be used to demonstrate the operation of your system.

3.3.1   User-Level Test Process

Write up to six user-level test processes to test your own OS. These test processes should run at unprivileged level and do not assume any kernel level data structures. These test processes only call the RTX APIs. The test processes should provide at least two and at most six test cases and finish testing within three minutes. The process id 1, 2, 3, 4, 5, and 6 are reserved for these processes.

Since the test processes have no knowledge of your detailed internal design, they only invoke the functions specified by the RTX API. The test processes can use the timer that is not used by the RTX for timing testing. We require the testing results to follow the following format and you output the results to the HyperTerminal (i.e. UART0):

```
Gid_test: START
Gid_test: test n OK
Gid_test: test m FAIL
Gid_test: x/N tests OK
Gid_test: y/N tests FAIL
Gid_test: END
```

For example, if you are group G099 and you have 3 testing cases in total. Two of the testing cases pass and one of the testing cases does not pass. The final testing results should be output to putty terminal as follows:

```
G099_test: START
G099_test: total 3 tests
G099_test: test 1 OK
G099_test: test 2 OK
G099_test: test 3 FAIL
G099_test: 2/3 tests OK
G099_test: 1/3 tests FAIL
G099_test: END
```

3.3.2   24 Hour Wall Clock Display Process

This process registers itself with the Keyboard Command Decoder process as the handler for the `%W` command.

- The `%WR` command will reset the current wall clock time to `00:00:00`, starts the clock running and causes display of the current wall clock time on the console CRT. The display will be updated every second.

- The `%WS hh:mm:ss` command sets the current wall clock time to `hh:mm:ss`, starts the clock running and causes display of the current wall clock time on the console CRT. The display will be updated every sec-ond.

- The `%WT` command will cause the wall clock display to be terminated.

3.3.3   Set Priority Command Process

This process registers itself with the Keyboard Command Decoder process as the handler for the `%C` command. The `%C` command has two parameters: `%C process_id new_priority` where `process_id` and `new_priority` are integers. This command changes the priority of the specified process, `process_id`, to `new_priority`. The change in priority level is immediate. It could also affect the target processs position on a ready queue or a blocked resource queue. The parameters must be verified to ensure a valid `process_id` and priority level is given. A `%C` command with illegal parameters will be ignored with an error message printed on the console.

3.3.4 Stress Test Processes

An important category of software tests are the stress tests. These tests seek to verify the behaviour of the system under heavy stress scenarios. One such scenario is depletion (or near depletion) of system resources. For the demonstration of this project, you will implement three processes whose behaviour is described below. The stress scenario being tested is depletion of memory blocks.

Process A:

```
p <- request_memory_block
register with Command Decoder as handler of %Z commands
loop forever
  p <- receive a message
  if the message(p) contains the %Z command then
    release_memory_block(p)
    exit the loop
  else
    release_memory_block(p)
  endif
endloop
num = 0
loop forever
  p <- request memory block to be used as a message envelope
  set message_type field of p to  count_report
  set msg_data[0] field of p to num
  send the message(p) to process B
  num = num + 1
  release_processor()
endloop
// note that Process A does not de-allocate
// any received envelopes in the second loop
```

Process B:

```
loop forever
  receive a message
  send the message to process C
endloop
```

Process C:

```
perform any needed initialization and create a local message queue
loop forever
  if (local message queue is empty) then
    p <- receive a message
  else
    p <- dequeue the first message from the local message queue
  endif

  if msg_type of p ==  count_report  then
    if msg_data[0] of p is evenly divisible by 20 then
      send "Process C" to CRT display using msg envelope p
      hibernate for 10 sec
    endif
  endif
  deallocate message envelope p
  release_processor()
endloop
```

The line "hibernate for 10 sec" is further expanded as:

```
q <- request_memory_block()
request a delayed_send for 10 sec delay with msg_type=wakeup10 using q
loop forever
  p <- receive a message //block and let other processes execute
  if (message_type of p == wakeup10) then
    exit this loop
  else
    put message (p) on the local message queue for later processing
  endif
endloop
```

Notes:

- Process C has a local message queue (distinct from the incoming message queue maintained by the RTX) onto which it enqueues (in FIFO order) messages which arrive while it hibernates. It processes these mes-sages later.

14

- For your own testing, set the priority levels for processes A, B and C to values which are most likely to cause memory block depletion in the RTX. During project demo, you may be asked to re-initialize your RTX with TA/instructor specified priorities for A, B, and C and vary the total number of message enve-lopes available.

3.4   Process ID Assignment

To facilitate the project evaluation, we enforce the process ID assignment rule listed in Table 3.1.

process ID	Process	Process ID	Process
0	Null	8	B
1	Test 1	9	C
2	Test 2	10	Set process priority process
3	Test 3	11	Wall clock display
4	Test 4	12	KCD
5	Test 5	13	CRT
6	Test 6	14	Timer i-process
7	A	15	UART i-process

Table 3.1: Required RTX Processes

Chapter 4

RTX Initialization

To make the RTX more generally applicable, the RTX will be configured at initialization as specified in the RTX Configuration Table. This table has three sections:

1. Memory configuration section: memory block size, number of memory blocks created;

2. System process section;

3. Application process section.

Chapter 3 (Required Processes) lists the processes to be created. Each entry contains the following data:

- process id,

- priority,

- stack size,

- start address, and

- for system processes, whether the process is an i-process.

All initializations must take place after the RTX execution starts.

Chapter 5

Deliverables and Demonstration

5.1 Deliverable

The project has four deliverables, where the first three deliverables are evaluated by in lab demonstrations. The deliverables are as follows:

5.1.1 RTX P1

This is the source code and documentation of a tiny kernel which provides memory management, processor management and process priority services. You need to implement

- APIs listed in Sections 2.1, 2.2 and 2.5;

- processes in Sections 3.1.1 and 3.3.1 [1]; and

- the corresponding part in Chapter 4.

5.1.2 RTX P2

This is the source code and documentation of a simplified version of the final RTX. On top of P1, you will

- add APIs listed in Sections 2.3 and 2.4;

---

[1]Note you need to write the six user testing processes to demonstrate that your implementation meets the requirements.

- finish implementing all processes as described in Sections 3.1, 3.2 and 3.3.2;

- enhance user test processes in Section 3.3.1 so that they test this version of the kernel; and

- implement the corresponding part in Chapter 4.

5.1.3   RTX P3

This is the final RTX source code to implement the specifications in Chapters 2, 3, and 4 based on the P1 and P2 implementations done previously. To be more specific, you will

- finish implementing processes as described in Sections 3.3.3 and 3.3.4;

- enhance user test processes in Section 3.3.1 so that they test the final version of the kernel; and

- implement the corresponding part in Chapter 4.

5.1.4   RTX P4

Program a second timer on the MCB1700 board to measure the speed of primitives. Write a final project report which include the following:

- RTX Project Software Design Description: This document describes the RTX design. It should include:

  - a structural description of the design (procedures and their interconnect; data structures; processes);
  - a functional description of all procedures (pseudo code; show all input/output parameters and globals);
  - implementation, testing and measurement plan (include responsibilities of individual team members).

  This part document should be kept reasonably small (no more than thirty pages, not including appendices). A design description should be shorter than the actual implementation.

- The measured times for the following primitives:

  - `send_message`,

- `receive_message` and
  - `request_memory_block`,

  and a check of the reasonableness of the values measured.

- A "lessons learned" summary (what you did do well, both technically and organizationally, and what you would do differently if you were to do it again). This summary should be brief (one to two pages).

The deliverable contains both the source code of the project with the timing code added in and the final report (in pdf format).

5.2 Demonstration

The first three deliverables of the project is evaluated by demonstration in the SE350 RTX Lab room.

5.2.1 Demo Reservation and Cancellation

- Use the Course Book System to reserve a time slot for a demo.

- You may cancel a demo reservation *48 hours before* the reserved deo time slot starts by using the course book system.

5.2.2 The Demo Policy

- During the demo of your project, your original submitted RTX implementation archive file will be retrieved and the demo will use those files in the archive. No substitutions are allowed.

- The demo is not some dry run to do additional debugging under "live" conditions. If minor bugs are discovered during the demo, depending on the complexity, you might be allowed to fix the bug, recompile and download and continue the demo. Under no circumstances will file replacements be allowed during the demo. "fixes" are basically limited to minor manual editing of a source file. Each demo for P1 and P2 is scheduled for a 30 minute slot maximum. The P3 demo slot is 50 minutes maximum.

- You are only allowed to demo once.

- ALL project group members are required to be presented during the demonstrations.

5.2.3   The Demo Procedure

P1 Demo Procedure

- **Basic Functionality Demo**
  You will need to demonstrate you have successfully completed the required APIs by using your own six test processes.

- **Source Code Spot Check**
  An evaluator will ask each group member implementation questions.

- **Contribution Check**
  Each group member will be asked what he/she has contributed to P1 implementation.

P2 Demo Procedure

- **Basic Functionality Demo**
  You will need to demonstrate you have successfully completed the required APIs by

  - displaying the wall clock display process using various `%W` commands and
  - showing the output of the three hot keys.

  Note you are responsible to program the six user test processes so that one process is blocked on memory and one process is blocked on receive to demonstrate the full functionality of the all the hot keys.

- **Source Code Spot Check**
  An evaluator will ask each group member implementation questions.

P3 Demo Procedure

- **Basic Functionality Demo**
  You will demonstrate the basic functionality of the RTX through commands line input. The evaluator will

  - start, observe and stop wall clock display by using various `%W` commands;
  - check the hot keys output; and
  - set priority of processes by using `%C` commands.

  test of wall clock display

- **Stress Test**
  Your RTX will be go through a stress testing demo by using processes A, B and C. The evaluator will

  - reinitialize RTX with `N` (`N=30`) envelopes, one combination of processes A,B and C priorities;

  - start test process activity by `%Z` command;

  - observe system operation (wall clock display as indicator), may also display trace buffer (is your have it implemented);

  The above will be repeated with two other combinations of processes A,B and C priorities.

- **Contribution Check**
  Each group member will be asked what he/she has contributed to the RTX (i.e. P1, P2 and P3) implementation.

Part II

Keil MCB1700 Quick Reference Guide

Chapter 6

Keil MCB1700 Hardware Environment

6.1   MCB1700 Board Overview

The Keil MCB1700 board is populated with *NXP LPC1768* Microcontroller. Figure 6.1 shows the important interface and hardware components of the MCB1700 board.

Figure 6.2 is the hardware block diagram that helps you to understand the MCB1700 board. Note that our lab will only use a small subset of the components which include the LPC1768 CPU, COM and Dual RS232.

The LPC1768 is a 32-bit ARM Cortex-M3 microcontroller for embedded applications requiring a high level of integration and low power dissipation. The LPC1768 operates at up to an 100 MHz CPU frequency. The peripheral complement of LPC1768 includes 512KB of on-chip flash memory, 64KB of on-chip SRAM and a variety of other on-chip peripherals. Among the on-chip peripherals, there are system control block, pin connect block, 4 UARTs and 4 general purpose timers, some of which will be used in your RTX course project. Figure 6.3 is the simplified LPC1768 block diagram [4], where the components to be used in your RTX project are circled with red. Note that this manual will only discuss the components that are relevant to the RTX course project. The LPC17xx User Manual is the complete reference for LPC1768 MCU.

6.2   Cortex-M3 Processor

The Cortex-M3 processor is the central processing unit (CPU) of the LPC1768 chip. The processor is a 32-bit microprocessor with a 32-bit data path, a 32-bit register bank, and

Figure 6.1: MCB1700 Board Components [1]

Figure 6.2: MCB1700 Board Block Diagram [1]
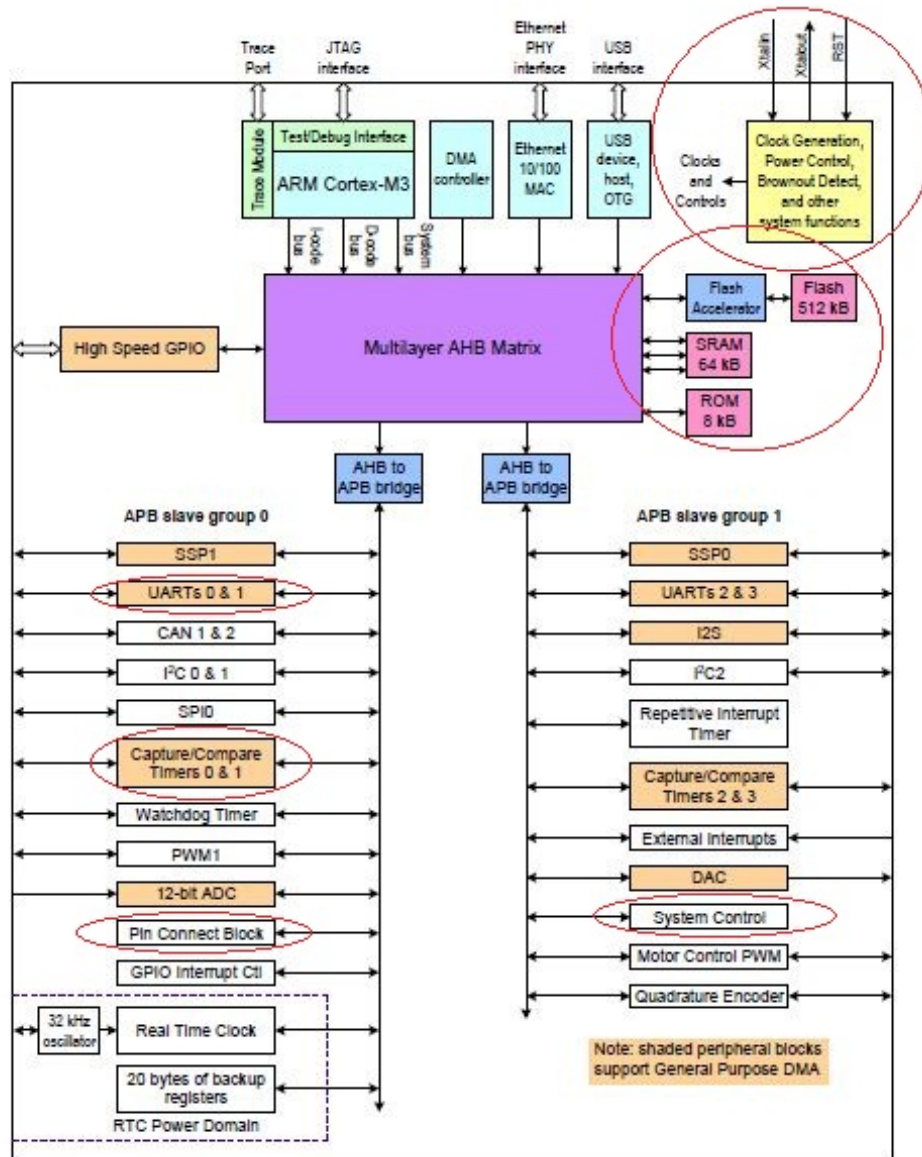
Figure 6.3: LPC1768 Block Diagram

Figure 6.4: Simplified Cortex-M3 Block Diagram[5]

32-bit memory interfaces. Figure 6.4 is the simplified block diagram of the Cortex-M3 processor [5]. The processor has private peripherals which are system control block, system timer, NVIC (Nested Vectored Interrupt Controller) and MPU (Memory Protection Unit). The MPU programming is not required in the course project. The processor includes a number of internal debugging components which provides debugging features such as breakpoints and watchpoints.

6.2.1 Registers

The processor core registers are shown in Figure 6.5. For detailed description of each register, Chapter 34 in [4] is the complete reference.

- R0-R12 are 32-bit general purpose registers for data operations. Some 16-bit Thumb instructions can only access the low registers (R0-R7).

- R13(SP) is the stack pointer alias for two banked registers shown as follows:

  - *Main Stack Pointer (MSP)*: This is the default stack pointer and also reset value. It is used by the OS kernel and exception handlers.

  - *Process Stack Pointer (PSP)*: This is used by user application code.

26

Figure 6.5: Cortex-M3 Registers[4]

On reset, the processor loads the MSP with the value from address 0x00000000. The lowest 2 bits of the stack pointers are always 0, which means they are always word aligned.

In Thread mode, when bit[1] of the CONTROL register is 0, MSP is used. When bit[1] of the CONTROL register is 1, PSP is used.

- R14(LR) is the link register. The return address of a subroutine is stored in the link register when the subroutine is called.

- R15(PC) is the program counter. It can be written to control the program flow.

- Special Registers are as follows:

  - Program Status registers (PSRs)
  - Interrupt Mask registers (PRIMASK, FAULTMASK, and BASEPRI)
  - Control register (CONTROL)

When at privilege level, all the registers are accessible. When at unprivileged (user) level, access to these registers are limited.

6.2.2 Processor mode and privilege levels

The Cortex-M3 processor supports two modes of operation, Thread mode and Handler mode.

- Thread mode is entered upon Reset and is used to execute application software.

- Handler mode is used to handle exceptions. The processor returns to Thread mode when it has finished exception handling.

Software execution has two access levels, Privileged level and Unprivileged (User) level.

- Privileged
  The software can use all instructions and has access to all resources. Your RTOS kernel functions are running in this mode.

- Unprivileged (User)
  The software has limited access to `MSR` and `MRS` instructions and cannot use the `CPS` instruction. There is no access to the system timer, NVIC , or system control block. The software might also have restricted access to memory or peripherals. User processes such as the wall clock process should run at this level.

When the processor is in Handler mode, it is at the privileged level. When the processor is in Thread mode, it can run at privileged or unprivileged (user) level. The bit[0] in CONTROL register determines the execution privilege level. Figure 6.6 illustrate the mode and privilege level of the processor.

Figure 6.6: Cortex-M3 Operating Mode and Privilege Level[5]

Note that only privileged software can write to the CONTROL register to change the privilege level for software execution in Thread mode. Unprivileged software can use the

SVC instruction to make a supervisor call to transfer control to privileged software. Another way to change between Privileged Thread mode and Unprivileged thread mode is to modify the EXC_RETURN value in the LR (R14) when returning from an exception. You probably want to use this mechanism for context switching.

6.2.3   Stacks

The processor uses a full descending stack. This means the stack pointer indicates the last stacked item on the stack memory. When the processor pushes a new item onto the stack, it decrements the stack pointer and then writes the item to the new memory location.

The processor implements two stacks, the *main stack* and the *process stack*. One of these two stacks is banked out depending on the stack in use. This means only one stack is visible at a time as R13. In Handler mode, the main stack is always used. The bit[1] in CONTROL register reads as zero and ignores writes in Handler mode. In Thread mode, the bit[1] setting in CONTROL register determines whether the main stack or the process stack is currently used. Table 6.1 summarizes the processor mode, execution privilege level, and stack use options.

Processor mode	Used to execute	Privilege level for software execution	CONTROL Bit[0]	Bit[1]	Stack used
Thread	Applications	Privileged	0	0	Main Stack
		Unprivileged	1	1	Process Stack
Handler	Exception handlers	Privileged	-	0	Main Stack

Table 6.1: Summary of processor mode, execution privilege level, and stack use options

6.3   Memory Map

The Cortex-M3 processor has a single fixed 4GB address space. Table 6.2 shows how this space is used on the LPC1768.

Note that the memory map is not continuous. For memory regions not shown in the table, they are reserved. When accessing reserved memory region, the processor's behavior is not defined. All the peripherals are memory-mapped and the LPC17xx.h file defines the data structure to access the memory-mapped peripherals in C.

Address Range	General Use	Address range details	Description
0x0000 0000 to 0x1FFF FFFF	On-chip non-volatile memory	0x0000 0000 − 0x0007 FFFF	512 KB flash memory
	On-chip SRAM	0x1000 0000 − 0x1000 7FFF	32 KB local SRAM
	Boot ROM	0x1FFF 0000 − 0x1FFF 1FFF	8 KB Boot ROM
0x2000 0000 to 0x3FFF FFFF	On-chip SRAM (typically used for peripheral data)	0x2007 C000 − 0x2007 FFFF	AHB SRAM - bank0 (16 KB)
		0x2008 0000 − 0x2008 3FFF	AHB SRAM - bank1 (16 KB)
	GPIO	0x2009 C000 − 0x2009 FFFF	GPIO
0x4000 0000 to 0x5FFF FFFF	APB Peripherals	0x4000 0000 − 0x4007 FFFF	APB0 Peripherals
		0x4008 0000 − 0x400F FFFF	APB1 Peripherals
	AHB peripherals	0x5000 0000 − 0x501F FFFF	DMA Controller, Ethernet interface, and USB interface
0xE000 0000 to 0xE00F FFFF	Cortex-M3 Private Peripheral Bus (PPB)	0xE000 0000 − 0xE00F FFFF	Cortex-M3 private registers(NVIC, MPU and SysTick Timer et. al.)

Table 6.2: LPC1768 Memory Map

6.4   Exceptions and Interrupts

The Cortex-M3 processor supports system exceptions and interrupts. The processor and the Nested Vectored Interrupt Controller (NVIC) prioritize and handle all exceptions. The processor uses *Handler mode* to handle all exceptions except for reset.

6.4.1   Vector Table

Exceptions are numbered 1-15 for system exceptions and 16 and above for external interrupt inputs. LPC1768 NVIC supports 35 vectored interrupts. Table 6.3 shows system exceptions and some frequently used interrupt sources. See Table 50 and Table 639 in [4] for the complete exceptions and interrupts sources. On system reset, the vector table is fixed at address 0x00000000. Privileged software can write to the VTOR (within the System Control Block) to relocate the vector table start address to a different memory location, in the range 0x00000080 to 0x3FFFFF80.

6.4.2   Exception Entry

Exception entry occurs when there is a pending exception with sufficient priority and either

- the processor is in Thread mode

Exception number	IRQ number	Vector address or offset	Exception type	Priority	C PreFix
1	-	0x00000004	Reset	-3, the highest	
2	-14	0x00000008	NMI	-2,	NMI_
3	-13	0x0000000C	Hard fault	-1	HardFault_
4	-12	0x00000010	Memory management fault	Configurable	MemManage_
⋮					
11	-5	0x0000002C	SVCall	Configurable	SVC_
⋮					
14	-2	0x00000038	PendSV	Configurable	PendSVC_
15	-1	0x0000003C	SysTick	Configurable	SysTick_
16	0	0x00000040	WDT	Configurable	WDT_IRQ
17	1	0x00000044	Timer0	Configurable	TIMER0_IRQ
18	2	0x00000048	Timer1	Configurable	TIMER1_IRQ
19	3	0x0000004C	Timer2	Configurable	TIMER2_IRQ
20	4	0x00000050	Timer3	Configurable	TIMER3_IRQ
21	5	0x00000054	UART0	Configurable	UART0_IRQ
22	6	0x00000058	UART1	Configurable	UART1_IRQ
23	7	0x0000005C	UART2	Configurable	UART2_IRQ
24	8	0x00000060	UART3	Configurable	UART3_IRQ
⋮					

Table 6.3: LPC1768 Exception and Interrupt Table

- the processor is in Handler mode and the new exception is of higher priority than the exception being handled, in which case the new exception preempts the original exception (This is the nested exception case which is not required in our RTOS lab).

When an exception takes place, the following happens

- Stacking

  When the processor invokes an exception (except for tail-chained or a late-arriving exception, which are not required in the RTOS lab), it automatically stores the following eight registers to the SP:

  - R0-R3, R12
  - PC (Program Counter)
  - PSR (Processor Status Register)
  - LR (Link Register, R14)

Figure 6.7 shows the exception stack frame. Note that by default the stack frame is aligned to double word address starting from Cortex-M3 revision 2. The alignment feature can be turned off by programming the STKALIGN bit in the System Control Block (SCB) Configuration Control Register (CCR) to 0. On exception entry, the processor uses bit[9] of the stacked PSR to indicate the stack alignment. On return from the exception, it uses this stacked bit to restore the correct stack alignment.


Figure 6.7: Cortex-M3 Exception Stack Frame [5]

- Vector Fetching
  While the data bus is busy stacking the registers, the instruction bus fetches the exception vector (the starting address of the exception handler) from the vector table. The stacking and vector fetch are performed on separate bus interfaces, hence they can be carried out at the same time.

- Register Updates
  After the stacking and vector fetch are completed, the exception vector will start to execute. On entry of the exception handler, the following registers will be updated as follows:

  – SP: The SP (MSP or PSP) will be updated to the new location during stacking. Stacking from the privileged/unprivileged thread to the first level of the exception handler uses the MSP/PSP. During the execution of exception handler routine, the MSP will be used when stack is accessed.

  – PSR: The IPSR will be updated to the new exception number

– PC: The PC will change to the vector handler when the vector fetch completes and starts fetching instructions from the exception vector.

– LR: The LR will be updated to a special value called EXC_RETURN. This indicates which stack pointer corresponds to the stack frame and what operation mode the processor was in before the exception entry occurred.

– Other NVIC registers: a number of other NVIC registers will be updated .For example the pending status of exception will be cleared and the active bit of the exception will be set.

6.4.3   EXC_RETURN Value

EXC_RETURN is the value loaded into the LR on exception entry. The exception mechanism relies on this value to detect when the processor has completed an exception handler. The EXC_RETURN bits $[31:4]$ is always set to 0xFFFFFFF by the processor. When this value is loaded into the PC, it indicates to the processor that the exception is complete and the processor initiates the exception return sequence. Table 6.4 describes the EXC_RETURN bit fields. Table 6.5 lists Cortex-M3 allowed EXC_RETURN values.

Bits	31:4	3	2	1	0
Description	0xFFFFFFF	Return mode (Thread/Handler)	Return stack	Reserved; must be 0	Process state (Thumb/ARM)

Table 6.4: EXC_RETURN bit fields [5]

Value	Description		
	Return Mode	Exception return gets state from	SP after return
0xFFFFFFF1	Handler	MSP	MSP
0xFFFFFFF9	Thread	MSP	MSP
0xFFFFFFFD	Thread	PSP	PSP

Table 6.5: EXC_RETURN Values on Cortex-M3

6.4.4   Exception Return

Exception return occurs when the processor is in Handler mode and executes one of the following instructions to load the EXC_RETURN value into the PC:

- a `POP` instruction that includes the PC. This is normally used when the `EXC_RETURN` in LR upon entering the exception is pushed onto the stack.

- a `BX` instruction with any register. This is normally used when LR contains the proper `EXC_RETURN` value before the exception return, then `BX LR` instruction will cause an exception return.

- a `LDR` or `LDM` instruction with the PC as the destination. This is another way to load PC with the `EXC_RETURN` value.

Note unlike the ColdFire processor which has the `RTE` as the special instruction for exception return, in Cortex-M3, a normal return instruction is used so that the whole interrupt handler can be implemented as a C subroutine.

When the exception return instruction is executed, the following exception return sequences happen:

- Unstacking: The registers (i.e. exception stack frame) pushed to the stack will be restored. The order of the POP will be the same as in stacking. The SP will also be changed back.

- NVIC register update: The active bit of the exception will be cleared. The pending bit will be set again if the external interrupt is still asserted, causing the processor to reenter the interrupt handler.

6.5   Data Types

The processor supports 32-bit words, 16-bit halfwords and 8-bit bytes. It supports 64-bit data transfer instructions. All data memory accesses are managed as little-endian.

Chapter 7

Keil Software Development Tools

The Keil MDK-ARM development tools are used for MCB1700 boards in our lab. The tools include

- μ`Vision4` IDE which combines the project manager, source code editor and program debugger into one environment;

- ARM compiler, assembler, linker and utilities;

- ULINK USB-JTAG Adapter which allows you to debug the embedded programs running on the board.

The MDK-Lite is the evaluation version and does not require a license. However it has a code size limit of 32KB, which is adequate for your course projects. The MDK-Lite version 4.60.0.0 [1] is installed on all lab computers. If you want to install the software on your own computer, Appendix A gives detailed instruction.

7.1   Creating an Application in μ`Vision4` IDE

To get started with the Keil IDE, the MDK-ARM Primer

$$http://www.keil.com/support/man/docs/gsac/$$

is a good place to start. We will walk you through the IDE by developing a simple HelloWorld application which displays Hello World through the UART0 that is connected to the lab PC. Note the HelloWorld example uses polling rather than interrupt.

---

[1]The latest version is 5.1.0.0 We have not fully tested the supplied sample code with this version. This manual is based on version 4.60.0.0.

7.1.1 Create a New Project

1. Create a folder named "HelloWorld" on your computer [2].

2. Copy the following files to "HelloWorld" folder:

   - manual_code\UART_polling\src\uart_polling.h
   - manual_code\UART_polling\src\uart_polling.c
   - manual_code\Startup\system_LPC17xx.c

3. Create a new μVision project by click

   - Project → New μVision Project (See Figure 7.1)

Figure 7.1: Keil IDE: Create a New Project

   - Choose NXP(Founded by Philips) → LPC1768 (See Figure 7.2(a) and Figure 7.2(b))
   - Answer "Yes" to copy the startup code (See Figure 7.3).

7.1.2 Managing Project Components

You just finished creating a new project. One the left side of the IDE is the Project window and expand all objects, you will see the default project setup as shown in Figure 7.4.

1. Rename the Target
   The "Target 1" is the default name of the project build target and you can rename it by clicking the target name to highlight it and then click the highlighted name to input a new target name, say "HelloWorld SIM"

---

[2]The folder path name should not contain spaces on Nexus computers.

(a) Choose NXP   (b) Choose LPC1768

Figure 7.2: Keil IDE: Choose MCU

Figure 7.3: Keil IDE: Copy Startup Code

2. Rename the Source Group
   The IDE allows you to group source files to different groups to better manage
   the source code. By default "Source Group 1" is created and the startup code
   "startup_LPC17xx.s" is put under this source group. You can rename the source
   group by clicking the source group name to highlight it and then click again to input
   a new name, say "Startup Code".

3. Add a New Source Group
   You can add new source groups to your project. Click "Project → Manage" →
   "Components, Environment, Books..." (See Figure 7.5 You can now add new source
   groups to the project. Let's add "System Code" and "Source Code" source groups
   to the project (See Figure 7.6. Your project will now look like Figure 7.7
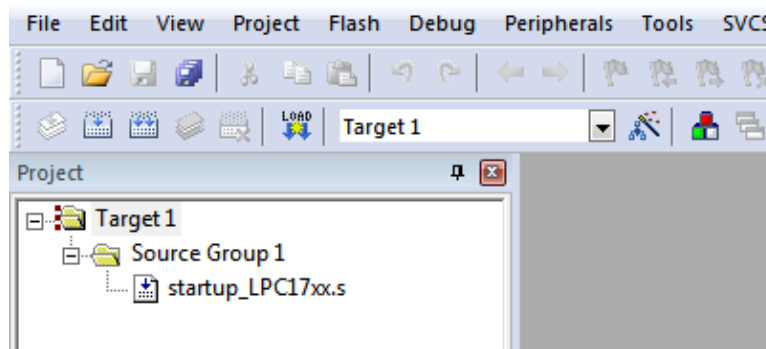
37

Figure 7.4: Keil IDE: A default new project

Figure 7.5: Keil IDE: Manage Project Components

Figure 7.6: Keil IDE: Manage Components Window

4. Add Source Code to a Source Group

   Now add "system_LPC17xx.c" to "System Code" group by double clicking the source group and choose the file from the file window. Double clicking the file name will add the file to the source group. Or you can select the file and click the "Add" button at the lower right corner of the window (See Figure 7.8).

   Similarly, add "uart_polling.c" to "Source Code" group. Your project will now look like Figure 7.9.

5. Create a new source file

   The project does not have a main function yet. We now create a new file by clicking the "New" button (See Figure 7.10). Before typing anything to the file, save the file
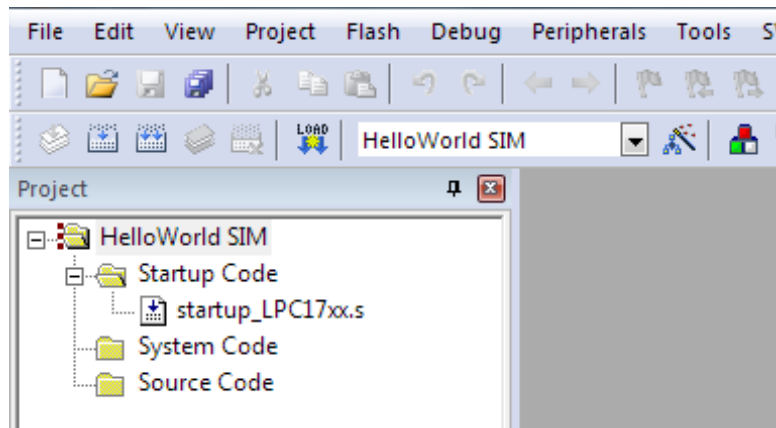
Figure 7.7: Keil IDE: Updated Project Profile
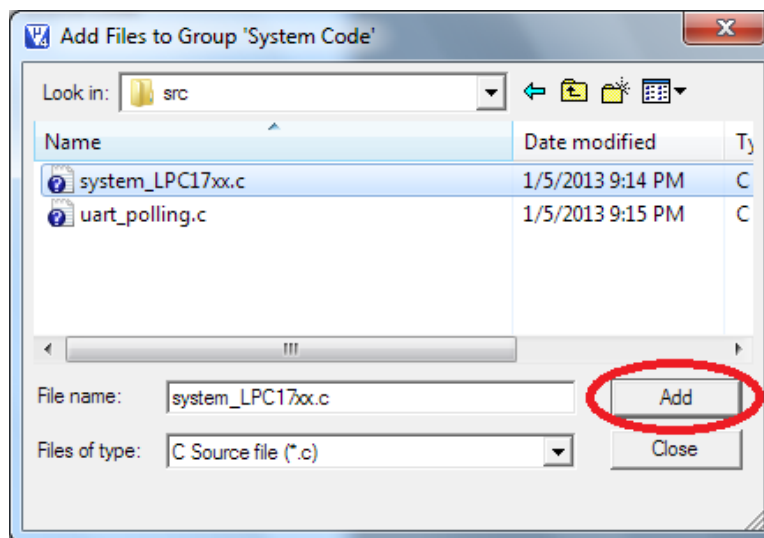
Figure 7.8: Keil IDE: Add Source File to Source Group

and name it "main.c". Put the following code to the main.c file:

```c
#include <LPC17xx.h>
#include "uart_polling.h"
int main() {
    SystemInit();
    uart0_init();
    uart0_put_string("Hello World!\n\r");
    return 0;
```
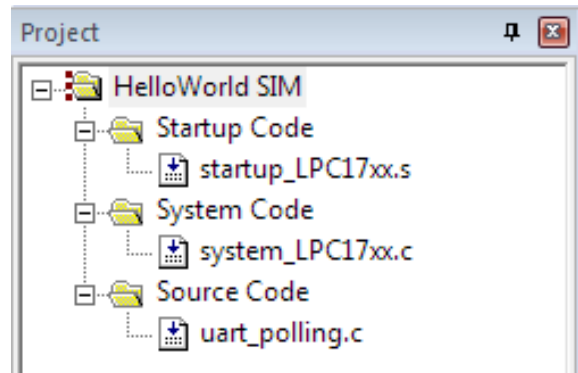
Figure 7.9: Keil IDE: Updated Project Profile

Figure 7.10: Keil IDE: Create New File

```
}
```

Then add main.c to the "Source Code" group. Your final project would look like the screen shot in Figure 7.11.

Figure 7.11: Keil IDE: Final Project Setting

7.1.3 Build and Download

To build the target, click the "Build" button (see Figure 7.12). If nothing is wrong, the

Figure 7.12: Keil IDE: Build Target

build output window at the bottom of the IDE will show a log similar like the one shown in Figure 7.13
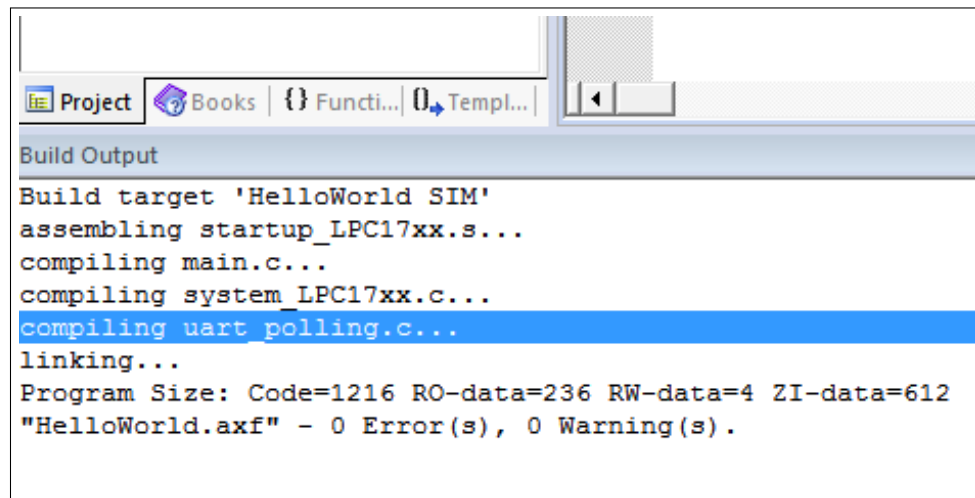
Figure 7.13: Keil IDE: Build Target

To download the code to the board, click the "Load" button (see Figure 7.14). The download is through the Ulink-Me.

You will need a terminal emulator such as PuTTY that talks directly to COM ports in order to see output of the serial port. Open up the PuTTY on your PC and choose COM1. An example PuTTY Serial configuration is shown in Figures 7.15(a) and 7.15(b). Press the Reset button on the board and you should see "Hello World!" displayed on PuTTY.

Figure 7.14: Keil IDE: Download Target to Flash
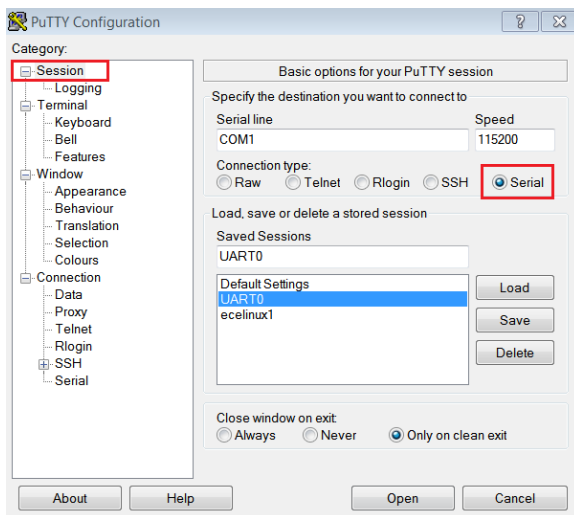
(a) PuTTY Session for Serial Port Communication
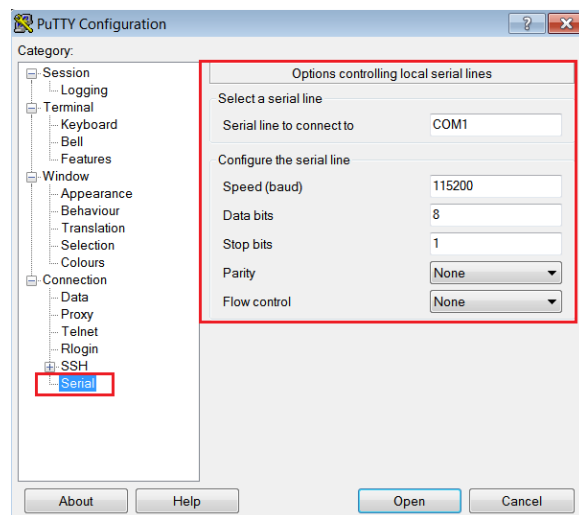
(b) PuTTY Serial Port Configuration

7.2 Debugging

You can use either the simulator within the IDE or the ULINK Cortex Debugger to debug your program. To start a debug session, click Debug→Start/Stop Debug Session from the IDE menu bar or press Ctrl+F5. Figure 7.15 shows the a typical debug session interface.

As any other GUI debugger, the IDE allows you to set up break points and step through your source code. It also shows the registers, which is very helpful for debugging low level code. Click View, Debug and Peripherals from the IDE menu bar and explore the functionality of the debugger.

7.2.1 Disabling CRP

In order to avoid stealing firmware , the LPC1768 provides Code Read Protection (CRP) that allows fine-grain control about which areas of the memory can be read. A detailed description is found in Section 32.6 of [4]. In essence if the Assembler Directive NO_CRP

42

Figure 7.15: Keil IDE: Debugging

is not present, the hardware is initialized to only make the firmware read-only (see Figure 7.16)

Figure 7.16: startup_LPC17xx.s excerpt

Since it is advisable to change values on the fly when debugging, the CRP should be disabled during prototyping. Open up the target option window and click the Asm tab. Put "NO_CRP" as shown in Figure 7.17

Figure 7.17: Keil IDE: Using Simulator for Debugging

7.2.2 Simulation

Most of the development normally is done under the simulation mode. The default setting of the project uses the simulator to debug as shown in the target option (see Figure 7.18 Instead of load the program to the board for execution, you can run the code using the

Figure 7.18: Keil IDE: Using Simulator for Debugging

debugger under simulation mode.

7.2.3 Configure In-Memory Execution Using ULINK Cortex Debugger

When you debug hardware related problems, you most likely will find the ULINK Cortex Debugger is helpful. You need to configure the debugger as shown in Figure 7.19.

Figure 7.19: Keil IDE: Using ULINK Cortex Debugger

The default image memory map setting is that the code is executed from the ROM (see Figure7.20(a). Since the ROM portion of the code needs to be flashed in order to

be executed on the board, this incurs wear-and-tear on the on-chip flash of the LPC1768. Since most attempts to write a functioning RTX will eventually require some more or less elaborate debugging, the flash memory might wear out quickly. Unlike the flash memory stick file systems where the wear is aimed to be uniformly distributed across the memory portion, this flash memory will get used over and over again in the same portion.

The ARM compiler can be configured to have a different starting address. We can create a RAM target where the code starting address is in RAM (see Figure 7.20(b). An initialization file `RAM.ini` (see Listing 7.1)is needed to do the proper setting of SP, PC and vector table offset register.

(a) Default Memory Setting

(b) In-Memory Execution Setting
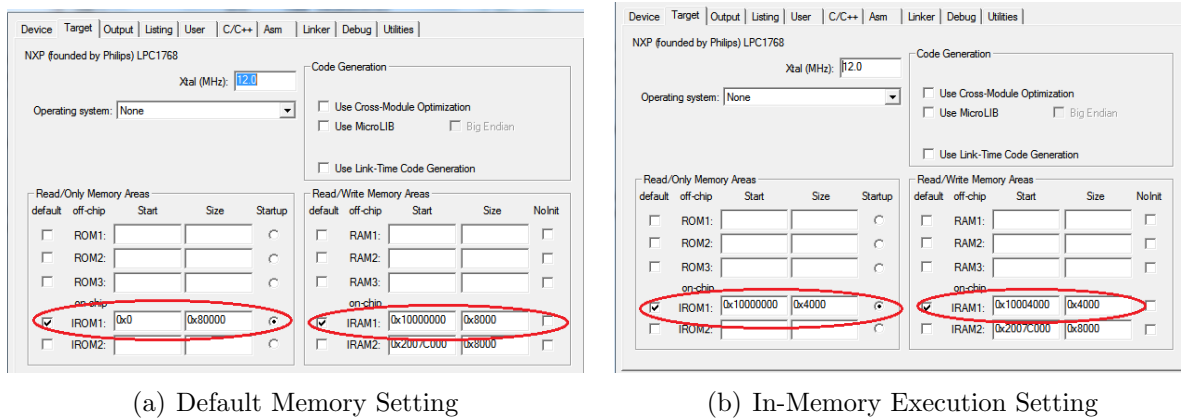
Figure 7.20: Keil IDE: Configure for In-Memory Execution

```
FUNC void Setup (void) {
  SP = _RDWORD(0x10000000);                     // Setup Stack Pointer
  PC = _RDWORD(0x10000004);                     // Setup Program Counter
  _WDWORD(0xE000ED08, 0x10000000);              // Setup Vector Table
      Offset Register
}
// You need to provide the path of the .axf file here
LOAD build\RAM\HelloWorld.axf INCREMENTAL      // Download
Setup();                                       // Setup for Running
g, main
```

Listing 7.1: The RAM.ini file

To download the code to the board, one should not use the download button. Instead, the debug button is used to initiate a debug session and the `RAM.ini` file will load the code to the board.

Chapter 8

Programming MCB1700

8.1 The Thumb-2 Instruction Set Architecture

The Cortex-M3 supports only the Thumb-2 (and traditional Thumb) instruction set. With support for both 16-bit and 32-bit instructions in the Thumb-2 instruction set, there is no need to switch the processor between Thumb state (16-bit instructions) and ARM state (32-bit instructions).

In the RTOS lab, you will need to program a little bit in the assembler language. We introduce a few assembly instructions that you most likely need to use in your project in this section.

The general formatting of the assembler code is as follows:

```
label
        opcode operand1, operand2, ... ; Comments
```

The `label` is optional. Normally the first operand is the destination of the operation (note `STR` is one exception).

Table 8.1 lists some assembly instructions that the RTX project may use. For complete instruction set reference, we refer the reader to Section 34.2 (ARM Cortex-M3 User Guide: Instruction Set) in [4].

8.2 ARM Architecture Procedure Call Standard (AAPCS)

The AAPCS (ARM Architecture Procedure Call Standard) defines how subroutines can be separately written, separately compiled, and separately assembled to work together. The

Mnemonic	Operands/Examples	Description
LDR	$Rt, [Rn, \#offset]$	Load Register with word
	LDR R1, [R0, #24]	Load word value from an memory address R0+24 into R1
LDM	$Rn\{!\}, reglist$	Load Multiple registers
	LDM R4, {R0 − R1}	Load word value from memory address R4 to R0, increment the address, load the value from the updated address to R1.
STR	$Rt, [Rn, \#offset]$	Store Register word
	STR R3, [R2, R6]	Store word in R3 to memory address R2+R6
	STR R1, [SP, #20]	Store word in R1 to memory address SP+20
MRS	$Rd, spec\_reg$	Move from special register to general register
	MRS R0, MSP	Read MSP into R0
	MRS R0, PSP	Read PSP into R0
MSR	$spec\_reg, Rm$	Move from general register to special register
	MSR MSP, R0	Write R0 to MSP
	MSR PSP, R0	Write R0 to PSP
PUSH	$reglist$	Push registers onto stack
	PUSH {R4 − R11, LR}	push in order of decreasing the register numbers
POP	$reglist$	Pop registers from stack
	POP {R4 − R11, PC}	pop in order of increasing the register numbers
BL	$label$	Branch with Link
	BL funC	Branch to address labeled by funC, return address stored in LR
BLX	Rm	Branch indirect with link
	BLX R12	Branch with link and exchange (Call) to an address stored in R12
BX	Rm	Branch indirect
	BX LR	Branch to address in LR, normally for function call return

Table 8.1: Assembler instruction examples

C compiler follows the AAPCS to generate the assembly code. Table 8.2 lists registers used by the AAPCS.

Register	Synonym	Special	Role in the procedure call standard
r15		PC	The Program Counter.
r14		LR	The Link Register.
r13		SP	The Stack Pointer (full descending stack).
r12		IP	The Intra-Procedure-call scratch register.
r11	v8		Variable-register 8.
r10	v7		Variable-register 7.
r9		v6	Platform register.
		SB	The meaning of this register is defined by platform standard.
		TR	
r8	v5		Variable-register 5.
r7	v4		Variable-register 4.
r6	v3		Variable-register 3.
r5	v2		Variable-register 2.
r4	v1		Variable-register 1.
r3	a4		argument / scratch register 4
r2	a3		argument / scratch register 3
r1	a2		argument / result / scratch register 2
r0	a1		argument / result / scratch register 1

Table 8.2: Core Registers and AAPCS Usage

Registers R0-R3 are used to pass parameters to a function and they are not preserved. The compiler does not generate assembler code to preserve the values of these registers. R0 is also used for return value of a function.

Registers R4-R11 are preserved by the called function. If the compiler generated assembler code uses registers in R4-R11, then the compiler generate assembler code to automatically push/pop the used registers in R4-R11 upon entering and exiting the function.

R12-R15 are special purpose registers. A function that has the `__svc_indirect` keyword makes the compiler put the first parameter in the function to R12 followed by an `SVC` instruction. R13 is the stack pointer (SP). R14 is the link register (LR), which normally is used to save the return address of a function. R15 is the program counter (PC).

Note that the exception stack frame automatically backs up R0-R3, R12, LR and PC together with the xPSR. This allows the possibility of writing the exception handler in purely C language without the need of having a small piece of assembly code to save/restore R0-R3, LR and PC upon entering/exiting an exception handler routine.

8.3 Cortex Microcontroller Software Interface Standard (CMSIS)

The Cortex Microcontroller Software Interface Standard (CMSIS) was developed by ARM. It provides a standardized access interface for embedded software products (see Figure 8.1). This improves software portability and re-usability. It enables software solution suppliers to develop products that can work seamlessly with device libraries from various silicon vendors [2].
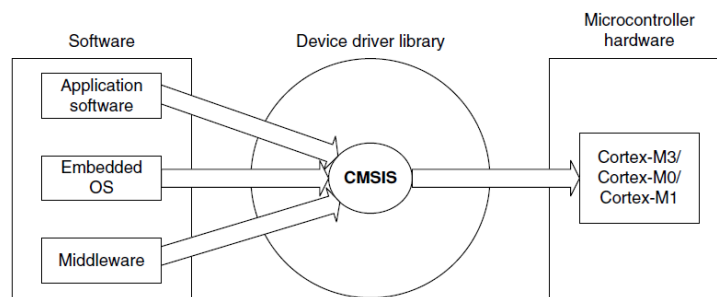
Figure 8.1: Role of CMSIS[5]

The CMSIS uses standardized methods to organize header files that makes it easy to learn new Cortex-M microcontroller products and improve software portability. With the `<device>.h` (e.g. `LPC17xx.h`) and system startup code files (e.g., `startup_LPC17xx.s`), your program has a common way to access

- **Cortex-M processor core registers** with standardized definitions for NVIC, SysTick, MPU registers, System Control Block registers , and their core access functions (see $core\_cm*.[ch]$ files).

- **system exceptions** with standardized exception number and handler names to allow RTOS and middleware components to utilize system exceptions without having compatibility issues.

- **intrinsic functions with standardized name** to produce instructions that cannot be generated by IEC/ISO C.

- **system initialization** by common methods for each MCU. Fore example, the standardized `SystemInit()` function to configure clock.

- **system clock frequency** with standardized variable named as `SystemFrequency` defined in the device driver.

- **vendor peripherals** with standardized C structure.

Figure 8.2: CMSIS Organization[2]

8.3.1 CMSIS files

The CMSIS is divided into multiple layers (See Figure 8.2). For each device, the MCU vendor provides a device header file `<device>.h` (e.g., `LPC17xx.h`) which pulls in additional header files required by the device driver library and the Core Peripheral Access Layer (see Figure 8.3).

By including the `<device>.h` (e.g., `LPC17xx.h`) file into your code file. The first step to initialize the system can be done by calling the CMSIS function as shown in Listing 8.1.

```
SystemInit(); // Initialize the MCU clock
```
Listing 8.1: CMSIS SystemInit()

The CMSIS compliant device drivers also contain a startup code (e.g., `startup_LPC17xx.s`), which include the vector table with standardized exception handler names (See Section 8.3.3.

8.3.2 Cortex-M Core Peripherals

We only introduce the NVIC programming in this section. The Nested Vectored Interrupt Controller (NVIC) can be accessed by using CMSIS functions (see Figure 8.4). As an

50

Figure 8.3: CMSIS Organization[2]

Figure 8.4: CMSIS NVIC Functions[2]

example, the following code enables the UART0 and TIMER0 interrupt

```
NVIC_EnableIRQ(UART0_IRQn); // UART0_IRQn is defined in LPC17xx.h
NVIC_EnableIRQ(TIMER0_IRQn); // TIMER0_IRQn is defined in LPC17xx.h
```

8.3.3   System Exceptions

Writing an exception handler becomes very easy. One just defines a function that takes no input parameter and returns void. The function takes the name of the standardized exception handler name as defined in the startup code (e.g., `startup_LPC17xx.s`). The following listing shows an example to write the UART0 interrupt handler entirely in C.

```
void UART0_Handler (void)
{
    // write your IRQ here
}
```

Another way is to use the embedded assembly code:

```
__asm void UART0_Handler(void)
{
    ; do some asm instructions here
    BL __cpp(a_c_function) ; a_c_function is a regular C function
    ; do some asm instructions here,
}
```

8.3.4   Intrinsic Functions

ANSI cannot directly access some Cortex-M3 instructions. The CMSIS provides intrinsic functions that can generate these instructions. The CMSIS also provides a number of functions for accessing the special registers using `MRS` and `MSR` instructions. The intrinsic functions are provided by the RealView Compiler. Table 8.3 lists some intrinsic functions that your RTOS project most likely will need to use. We refer the reader to Tables 613 and 614 one page 650 in Section 34.2.2 of [4] for the complete list of intrinsic functions.

8.3.5   Vendor Peripherals

All vendor peripherals are organized as C structure in the `<device>.h` file (e.g., `LPC17xx.h`). For example, to read a character received in the RBR of UART0, we can use the following code.

```
unsigned char ch;
ch = LPC_UART0->RBR; // read UART0 RBR and save it in ch
```

Instruction		CMSIS Intrinsic Function
CPSIE I		void __enable_irq(void)
CPSID I		void __disable_irq(void)
Special Register	Access	CMSIS Function
CONTROL	Read	uinit32_t __get_CONTROL(void)
	Write	void __set_CONTROL(uint32_t value)
MSP	Read	uinit32_t __get_MSP(void)
	Write	void __set_MSP(uint32_t value)
PSP	Read	uinit32_t __get_PSP(void)
	Write	void __set_PSP(uint32_t value)

Table 8.3: CMSIS intrinsic functions defined in `core_cmFunc.h`

8.4   Accessing C Symbols from Assembly

Only embedded assembly is support in Cortex-M3. To write an embedded assembly function, you need to use the `__asm` keyword. For example the the function "`embedded_asm_function`" in Listing 8.2 is an embedded assembly function. You can only put assembly instructions inside this function. Note that inline assembly is not supported in Cortex-M3.

The `__cpp` keyword allows one to access C compile-time constant expressions, including the addresses of data or functions with external linkage, from the assembly code. The expression inside the `__cpp` can be one of the followings:

- A global variable defined in C

```
#define U32 unsigned int

typedef struct pcb {
    struct pcb *mp_next;
    U32 *mp_sp; //4 bytes offset from
                //the starting address of the structure
    //other variables...
} PCB;

#define SP_OFFSET 4

PCB g_pcb;
U32 g_var;
```

53

```
__asm embedded_asm_function(void) {
    LDR R3, =__cpp(&g_pcb) ; load R3 with the address of g_pcb
    LDM R3, {R1, R2}        ; load R1 with g_pcb.mp_next
                            ; load R2 with g_pcb.mp_sp
    LDR R4, =__cpp(g_var)  ; load R4 with the value of g_var
    STR R4, [R3, #SP_OFFSET] ; write R4 value to g_pcb.mp_sp
}
```

Listing 8.2: Example of accessing global variable from assembly

- A C function

```
extern void a_c_function(void);
...
__asm embedded_asm_function(void) {
    ;......
    BL __cpp(a_c_function) ; a_c_function is regular C function
    ;......
}
```

- A constant expression in the range of $0 - 255$ defined in C.

```
unsigned char const g_flag;

__asm embedded_asm_function(void) {
    ;......
    MOV R4, #_cpp(g_flag) ; load g_flag value to R4
    ;......
}
```

Note the MOV instruction only applies to immediate constant value in the range of $0 - 255$.

You can also use the IMPORT directive to import a C symbol in the embedded assembly function and then start to use the imported symbol just as a regular assembly symbol. For example

```
void a_c_function (void) {
   // do something
}
```

54

```
__asm embedded_asm_add(void) {
    IMPORT a_c_function ; a_c_function is a regular C function
    BL a_c_function   ; branch with link to a_c_function
}
```

Names in the `__cpp` expression are looked up in the C context of the `__asm` function. Any names in the result of the `__cpp` expression are mangled as required and automatically have `IMPORT` statements generated from them.

8.5   UART Programming

To program a UART on MCB1700 board, one first needs to configure the UART by following the steps listed in Section 15.1 in [4] (referred as LPC17xx_UM in the sample code comments). Listings 8.3, 8.4 and 8.5 give one sample implementation of programming UART0 interrupts.

```
/**
 * @brief: UART defines
 * @file: uart_def.h
 * @author: Yiqing Huang
 * @date: 2014/02/08
 */

#ifndef UART_DEF_H_
#define UART_DEF_H_

/* The following macros are from NXP uart.h */
#define IER_RBR  0x01
#define IER_THRE 0x02
#define IER_RLS  0x04

#define IIR_PEND 0x01
#define IIR_RLS  0x03
#define IIR_RDA  0x02
#define IIR_CTI  0x06
#define IIR_THRE 0x01

#define LSR_RDR  0x01
```

```c
#define LSR_OE   0x02
#define LSR_PE   0x04
#define LSR_FE   0x08
#define LSR_BI   0x10
#define LSR_THRE 0x20
#define LSR_TEMT 0x40
#define LSR_RXFE 0x80

#define BUFSIZE   0x40
/* end of NXP uart.h file reference */


/* convenient macro for bit operation */
#define BIT(X)  ( 1 << X )

/*
   8 bits, no Parity, 1 Stop bit

   0x83 = 1000 0011 = 1 0 00 0 0 11
   LCR[7] =1  enable Divisor Latch Access Bit DLAB
   LCR[6] =0  disable break transmission
   LCR[5:4]=00 odd parity
   LCR[3] =0  no parity
   LCR[2] =0  1 stop bit
   LCR[1:0]=11 8-bit char len
   See table 279, pg306 LPC17xx_UM
*/
#define UART_8N1 0x83

#ifndef NULL
#define NULL 0
#endif



#endif /* !UART_DEF_H_ */
```

Listing 8.3: UART0 IRQ Sample Code uart_def.h

```c
/**
 * @brief: uart.h
 * @author: Yiqing Huang
```

```c
 * @date: 2014/02/08
 */

#ifndef UART_IRQ_H_
#define UART_IRQ_H_

/* typedefs */
#include <stdint.h>
#include "uart_def.h"

/* The following macros are from NXP uart.h */
/*
#define IER_RBR  0x01
#define IER_THRE 0x02
#define IER_RLS  0x04

#define IIR_PEND 0x01
#define IIR_RLS  0x03
#define IIR_RDA  0x02
#define IIR_CTI  0x06
#define IIR_THRE 0x01

#define LSR_RDR  0x01
#define LSR_OE   0x02
#define LSR_PE   0x04
#define LSR_FE   0x08
#define LSR_BI   0x10
#define LSR_THRE 0x20
#define LSR_TEMT 0x40
#define LSR_RXFE 0x80

#define BUFSIZE  0x40
*/
/* end of NXP uart.h file reference */


/* convenient macro for bit operation */
//#define BIT(X)  ( 1 << X )

/*
```

```
  8 bits, no Parity, 1 Stop bit

  0x83 = 1000 0011 = 1 0 00 0 0 11
  LCR[7] =1  enable Divisor Latch Access Bit DLAB
  LCR[6] =0  disable break transmission
  LCR[5:4]=00 odd parity
  LCR[3] =0  no parity
  LCR[2] =0  1 stop bit
  LCR[1:0]=11 8-bit char len
  See table 279, pg306 LPC17xx_UM
*/
//#define UART_8N1 0x83


#define uart0_irq_init() uart_irq_init(0)
#define uart1_irq_init() uart_irq_init(1)

/* initialize the n_uart to use interrupt */
int uart_irq_init(int n_uart);

#endif /* ! UART_IRQ_H_ */
```

Listing 8.4: UART0 IRQ Sample Code uart.h

```
/**
 * @brief: uart_irq.c
 * @author: NXP Semiconductors
 * @author: Y. Huang
 * @date: 2014/02/08
 */

#include <LPC17xx.h>
#include "uart.h"
#include "uart_polling.h"
#ifdef DEBUG_0
#include "printf.h"
#endif


uint8_t g_buffer[]= "You Typed a Q\n\r";
uint8_t *gp_buffer = g_buffer;
```

```c
uint8_t g_send_char = 0;
uint8_t g_char_in;
uint8_t g_char_out;

/**
 * @brief: initialize the n_uart
 * NOTES: It only supports UART0. It can be easily extended to support
     UART1 IRQ.
 * The step number in the comments matches the item number in Section 14.1
     on pg 298
 * of LPC17xx_UM
 */
int uart_irq_init(int n_uart) {

  LPC_UART_TypeDef *pUart;

  if ( n_uart ==0 ) {
    /*
    Steps 1 & 2: system control configuration.
    Under CMSIS, system_LPC17xx.c does these two steps


    -----------------------------------------------------
    Step 1: Power control configuration.
            See table 46 pg63 in LPC17xx_UM
    -----------------------------------------------------
    Enable UART0 power, this is the default setting
    done in system_LPC17xx.c under CMSIS.
    Enclose the code for your refrence
    //LPC_SC->PCONP |= BIT(3);


    -----------------------------------------------------
    Step2: Select the clock source.
          Default PCLK=CCLK/4 , where CCLK = 100MHZ.
          See tables 40 & 42 on pg56-57 in LPC17xx_UM.
    -----------------------------------------------------
    Check the PLL0 configuration to see how XTAL=12.0MHZ
    gets to CCLK=100MHZin system_LPC17xx.c file.
    PCLK = CCLK/4, default setting after reset.
    Enclose the code for your reference
    //LPC_SC->PCLKSEL0 &= ~(BIT(7)|BIT(6));
```

```
  --------------------------------------------------------
  Step 5: Pin Ctrl Block configuration for TXD and RXD
          See Table 79 on pg108 in LPC17xx_UM.
  --------------------------------------------------------
  Note this is done before Steps3-4 for coding purpose.
  */

  /* Pin P0.2 used as TXD0 (Com0) */
  LPC_PINCON->PINSEL0 |= (1 << 4);

  /* Pin P0.3 used as RXD0 (Com0) */
  LPC_PINCON->PINSEL0 |= (1 << 6);

  pUart = (LPC_UART_TypeDef *) LPC_UART0;

} else if ( n_uart == 1) {

  /* see Table 79 on pg108 in LPC17xx_UM */
  /* Pin P2.0 used as TXD1 (Com1) */
  LPC_PINCON->PINSEL4 |= (2 << 0);

  /* Pin P2.1 used as RXD1 (Com1) */
  LPC_PINCON->PINSEL4 |= (2 << 2);

  pUart = (LPC_UART_TypeDef *) LPC_UART1;

} else {
  return 1; /* not supported yet */
}

/*
--------------------------------------------------------
Step 3: Transmission Configuration.
        See section 14.4.12.1 pg313-315 in LPC17xx_UM
        for baud rate calculation.
--------------------------------------------------------
      */

/* Step 3a: DLAB=1, 8N1 */
```

```
pUart->LCR = UART_8N1; /* see uart.h file */

/* Step 3b: 115200 baud rate @ 25.0 MHZ PCLK */
pUart->DLM = 0; /* see table 274, pg302 in LPC17xx_UM */
pUart->DLL = 9; /* see table 273, pg302 in LPC17xx_UM */

/* FR = 1.507 ~ 1/2, DivAddVal = 1, MulVal = 2
   FR = 1.507 = 25MHZ/(16*9*115200)
   see table 285 on pg312 in LPC_17xxUM
*/
pUart->FDR = 0x21;




/*
--------------------------------------------------------
Step 4: FIFO setup.
        see table 278 on pg305 in LPC17xx_UM
--------------------------------------------------------
      enable Rx and Tx FIFOs, clear Rx and Tx FIFOs
Trigger level 0 (1 char per interrupt)
*/

pUart->FCR = 0x07;

/* Step 5 was done between step 2 and step 4 a few lines above */

/*
--------------------------------------------------------
Step 6 Interrupt setting and enabling
--------------------------------------------------------
*/
/* Step 6a:
   Enable interrupt bit(s) wihtin the specific peripheral register.
         Interrupt Sources Setting: RBR, THRE or RX Line Stats
   See Table 50 on pg73 in LPC17xx_UM for all possible UART0 interrupt
      sources
   See Table 275 on pg 302 in LPC17xx_UM for IER setting
*/
/* disable the Divisior Latch Access Bit DLAB=0 */
```

```c
  pUart->LCR &= ~(BIT(7));

  //pUart->IER = IER_RBR | IER_THRE | IER_RLS;
  pUart->IER = IER_RBR | IER_RLS;

  /* Step 6b: enable the UART interrupt from the system level */

  if ( n_uart == 0 ) {
    NVIC_EnableIRQ(UART0_IRQn); /* CMSIS function */
  } else if ( n_uart == 1 ) {
    NVIC_EnableIRQ(UART1_IRQn); /* CMSIS function */
  } else {
    return 1; /* not supported yet */
  }
  pUart->THR = '\0';
  return 0;
}


/**
 * @brief: use CMSIS ISR for UART0 IRQ Handler
 * NOTE: This example shows how to save/restore all registers rather than
     just
 *       those backed up by the exception stack frame. We add extra
 *       push and pop instructions in the assembly routine.
 *       The actual c_UART0_IRQHandler does the rest of irq handling
 */
__asm void UART0_IRQHandler(void)
{
  PRESERVE8
  IMPORT c_UART0_IRQHandler
  PUSH{r4-r11, lr}
  BL c_UART0_IRQHandler
  POP{r4-r11, pc}
}
/**
 * @brief: c UART0 IRQ Handler
 */
void c_UART0_IRQHandler(void)
{
```

```c
  uint8_t IIR_IntId;    // Interrupt ID from IIR
  LPC_UART_TypeDef *pUart = (LPC_UART_TypeDef *)LPC_UART0;

#ifdef DEBUG_0
  uart1_put_string("Entering c_UART0_IRQHandler\n\r");
#endif // DEBUG_0

  /* Reading IIR automatically acknowledges the interrupt */
  IIR_IntId = (pUart->IIR) >> 1 ; // skip pending bit in IIR
  if (IIR_IntId & IIR_RDA) { // Receive Data Avaialbe
    /* read UART. Read RBR will clear the interrupt */
    g_char_in = pUart->RBR;
#ifdef DEBUG_0
    uart1_put_string("Reading a char = ");
    uart1_put_char(g_char_in);
    uart1_put_string("\n\r");
#endif // DEBUG_0

    g_buffer[12] = g_char_in; // nasty hack
    g_send_char = 1;
  } else if (IIR_IntId & IIR_THRE) {
  /* THRE Interrupt, transmit holding register becomes empty */

    if (*gp_buffer != '\0' ) {
      g_char_out = *gp_buffer;
#ifdef DEBUG_0
      //uart1_put_string("Writing a char = ");
      //uart1_put_char(g_char_out);
      //uart1_put_string("\n\r");

      // you could use the printf instead
      printf("Writing a char = %c \n\r", g_char_out);
#endif // DEBUG_0
      pUart->THR = g_char_out;
      gp_buffer++;
    } else {
#ifdef DEBUG_0
      uart1_put_string("Finish writing. Turning off IER_THRE\n\r");
#endif // DEBUG_0
      pUart->IER ^= IER_THRE; // toggle the IER_THRE bit
```

```
        pUart->THR = '\0';
        g_send_char = 0;
        gp_buffer = g_buffer;
    }

  } else { /* not implemented yet */
#ifdef DEBUG_0
      uart1_put_string("Should not get here!\n\r");
#endif // DEBUG_0
    return;
  }
}
```

Listing 8.5: UART0 IRQ Sample Code uart_irq.c

Listings 8.6 and 8.7 give one sample implementation of programming UART0 by polling.

```
/**
 * @brief: uart_polling.h
 * @author: Yiqing Huang
 * @date:   2014/01/05
 */

#ifndef UART_POLLING_H_
#define UART_POLLING_H_

#include <stdint.h> /* typedefs */
#include "uart_def.h"

#define uart0_init()     uart_init(0)
#define uart0_get_char() uart_get_char(0)
#define uart0_put_char(c) uart_put_char(0,c)
#define uart0_put_string(s) uart_put_string(0,s)

#define uart1_init()     uart_init(1)
#define uart1_get_char() uart_get_char(1)
#define uart1_put_char(c) uart_put_char(1,c)
#define uart1_put_string(s) uart_put_string(1,s)

int uart_init(int n_uart);   /* initialize the n_uart    */
int uart_get_char(int n_uart); /* read a char from the n_uart */
```

64

```
int uart_put_char(int n_uart, unsigned char c); /* write a char to n_uart
    */
int uart_put_string(int n_uart, unsigned char *s);/* write a string to
    n_uart */
void putc(void *p, char c);  /* call back function for printf, use uart1
    */


#endif /* ! UART_POLLING_H_ */
```

Listing 8.6: UART0 IRQ Sample Code uart_polling.h

```
/**
 * @brief: uart_polling.c, polling UART to send and receive data
 * @author: Yiqing Huang
 * @date: 2014/01/05
 * NOTE: the code only handles UART0 for now.
 */

#include <LPC17xx.h>
#include "uart_polling.h"


/**
 * @brief: initialize the n_uart
 * NOTES: only tested uart0 so far, but can be easily extended to other
    uarts.
 *        it should work with uart1, but no testing was done.
 */
int uart_init(int n_uart) {

  LPC_UART_TypeDef *pUart; /* ptr to memory mapped device UART, check  */
                        /* LPC17xx.h for UART register C structure
                              overlay */

  if (n_uart == 0 ) {
    /*
    Step 1: system control configuration

    step 1a: power control configuration, table 46 pg63
    enable UART0 power, this is the default setting
    also already done in system_LPC17xx.c
```

```
    enclose the code below for reference
    LPC_SC->PCONP |= BIT(3);


    step 1b: select the clock source, default PCLK=CCLK/4 , where CCLK =
        100MHZ.
    tables 40 and 42 on pg56 and pg57
    Check the PLL0 configuration to see how XTAL=12.0MHZ gets to CCLK=100
        MHZ
    in system_LPC17xx.c file
    enclose code below for reference
    LPC_SC->PCLKSEL0 &= ~(BIT(7)|BIT(6)); // PCLK = CCLK/4, default setting
        after reset

    Step 2: Pin Ctrl Block configuration for TXD and RXD
    Listed as item #5 in LPC_17xxum UART0/2/3 manual pag298
    */
    LPC_PINCON->PINSEL0 |= (1 << 4); /* Pin P0.2 used as TXD0 (Com0) */
    LPC_PINCON->PINSEL0 |= (1 << 6); /* Pin P0.3 used as RXD0 (Com0) */

    pUart = (LPC_UART_TypeDef *) LPC_UART0;

} else if (n_uart == 1) {
    LPC_PINCON->PINSEL4 |= (2 << 0); /* Pin P2.0 used as TXD1 (Com1) */
    LPC_PINCON->PINSEL4 |= (2 << 2); /* Pin P2.1 used as RXD1 (Com1) */

    pUart = (LPC_UART_TypeDef *) LPC_UART1;

} else {
    return -1;  /* not supported yet */
}

/* Step 3: Transmission Configuration */

/* step 3a: DLAB=1, 8N1 */
pUart->LCR   = UART_8N1;

/* step 3b: 115200 baud rate @ 25.0 MHZ PCLK */
pUart->DLM = 0;
pUart->DLL = 9;
```

```c
  pUart->FDR = 0x21;        /* FR = 1.507 ~ 1/2, DivAddVal = 1, MulVal = 2 */
                            /* FR = 1.507 = 25MHZ/(16*9*115200)          */
  pUart->LCR &= ~(BIT(7)); /* disable the Divisior Latch Access Bit DLAB=0
       */

  return 0;
}


/**
 * @brief: read a char from the n_uart, blocking read
 */

int uart_get_char(int n_uart)
{
  LPC_UART_TypeDef *pUart;

  if (n_uart == 0) {
    pUart = (LPC_UART_TypeDef *) LPC_UART0;
  } else if (n_uart == 1) {
    pUart = (LPC_UART_TypeDef *) LPC_UART1;
  } else {
    return -1; /* UART2,3 not supported yet */
  }

  /* polling the LSR RDR (Receiver Data Ready) bit to wait it is not empty
       */
  while (!(pUart->LSR & LSR_RDR));
  return (pUart->RBR);
}

/**
 * @brief: write a char c to the n_uart
 */

int uart_put_char(int n_uart, unsigned char c)
{
  LPC_UART_TypeDef *pUart;

  if (n_uart == 0) {
```

```c
    pUart = (LPC_UART_TypeDef *)LPC_UART0;
  } else if (n_uart == 1) {
    pUart = (LPC_UART_TypeDef *)LPC_UART1;
  } else {
    return -1; // UART2,3 not supported
  }

  /* polling LSR THRE bit to wait it is empty */
  while (!(pUart->LSR & LSR_THRE));
  return (pUart->THR = c); /* write c to the THR */
}

/**
 * @brief write a string to UART
 */
int uart_put_string(int n_uart, unsigned char *s)
{
  if (n_uart >1 ) return -1; /* only uart0, 1 are supported for now  */
  while (*s !=0) {            /* loop through each char in the string */
    uart_put_char(n_uart, *s++);/* print the char, then ptr increments */
  }
  return 0;
}

/**
 * @brief call back function for printf
 * NOTE: first paramter p is not used for now. UART1 used.
 */
void putc(void *p, char c)
{
  if ( p != NULL ) {
    uart1_put_string("putc: first parameter needs to be NULL");
  } else {
      uart1_put_char(c);
  }
}
```

Listing 8.7: UART0 IRQ Sample Code uart_polling.c

8.6 Timer Programming

To program a TIMER on MCB1700 board, one first needs to configure the TIMER by following the steps listed in Section 21.1 in [4]. Listings 8.8 and 8.9 give one sample implementation of programming TIMER0 interrupts. The timer interrupt fires every one millisecond.

```
/**
 * @brief timer.h - Timer header file
 * @author Y. Huang
 * @date 2013/02/12
 */
#ifndef _TIMER_H_
#define _TIMER_H_

extern uint32_t timer_init ( uint8_t n_timer ); /* initialize timer
    n_timer */

#endif /* ! _TIMER_H_ */
```
Listing 8.8: Timer0 IRQ Sample Code timer.h

```
/**
 * @brief timer.c - Timer example code. Tiemr IRQ is invoked every 1ms
 * @author T. Reidemeister
 * @author Y. Huang
 * @author NXP Semiconductors
 * @date 2012/02/12
 */

#include <LPC17xx.h>
#include "timer.h"

#define BIT(X) (1<<X)

volatile uint32_t g_timer_count = 0; // increment every 1 ms

/**
 * @brief: initialize timer. Only timer 0 is supported
 */
uint32_t timer_init(uint8_t n_timer)
```

```c
{
  LPC_TIM_TypeDef *pTimer;
  if (n_timer == 0) {
    /*
    Steps 1 & 2: system control configuration.
    Under CMSIS, system_LPC17xx.c does these two steps


    -------------------------------------------------------
    Step 1: Power control configuration.
            See table 46 pg63 in LPC17xx_UM
    -------------------------------------------------------
    Enable UART0 power, this is the default setting
    done in system_LPC17xx.c under CMSIS.
    Enclose the code for your refrence
    //LPC_SC->PCONP |= BIT(1);


    -------------------------------------------------------
    Step2: Select the clock source,
           default PCLK=CCLK/4 , where CCLK = 100MHZ.
           See tables 40 & 42 on pg56-57 in LPC17xx_UM.
    -------------------------------------------------------
    Check the PLL0 configuration to see how XTAL=12.0MHZ
    gets to CCLK=100MHZ in system_LPC17xx.c file.
    PCLK = CCLK/4, default setting in system_LPC17xx.c.
    Enclose the code for your reference
    //LPC_SC->PCLKSEL0 &= ~(BIT(3)|BIT(2));


    -------------------------------------------------------
    Step 3: Pin Ctrl Block configuration.
           Optional, not used in this example
           See Table 82 on pg110 in LPC17xx_UM
    -------------------------------------------------------
    */
    pTimer = (LPC_TIM_TypeDef *) LPC_TIM0;

  } else { /* other timer not supported yet */
    return 1;
  }

  /*
```

```
   --------------------------------------------------------
   Step 4: Interrupts configuration
   --------------------------------------------------------
   */

   /* Step 4.1: Prescale Register PR setting
      CCLK = 100 MHZ, PCLK = CCLK/4 = 25 MHZ
      2*(12499 + 1)*(1/25) * 10^(-6) s = 10^(-3) s = 1 ms
      TC (Timer Counter) toggles b/w 0 and 1 every 12500 PCLKs
      see MR setting below
   */
   pTimer->PR = 12499;

   /* Step 4.2: MR setting, see section 21.6.7 on pg496 of LPC17xx_UM. */
   pTimer->MR0 = 1;

   /* Step 4.3: MCR setting, see table 429 on pg496 of LPC17xx_UM.
      Interrupt on MR0: when MR0 mathches the value in the TC,
                        generate an interrupt.
      Reset on MR0: Reset TC if MR0 mathches it.
   */
   pTimer->MCR = BIT(0) | BIT(1);

   g_timer_count = 0;

   /* Step 4.4: CSMSIS enable timer0 IRQ */
   NVIC_EnableIRQ(TIMER0_IRQn);

   /* Step 4.5: Enable the TCR. See table 427 on pg494 of LPC17xx_UM. */
   pTimer->TCR = 1;

   return 0;
}

/**
 * @brief: use CMSIS ISR for TIMER0 IRQ Handler
 * NOTE: This example shows how to save/restore all registers rather than
      just
 *       those backed up by the exception stack frame. We add extra
 *       push and pop instructions in the assembly routine.
```

```
 *       The actual c_TIMER0_IRQHandler does the rest of irq handling
 */
__asm void TIMER0_IRQHandler(void)
{
  PRESERVE8
  IMPORT c_TIMER0_IRQHandler
  PUSH{r4-r11, lr}
  BL c_TIMER0_IRQHandler
  POP{r4-r11, pc}
}
/**
 * @brief: c TIMER0 IRQ Handler
 */
void c_TIMER0_IRQHandler(void)
{
  /* ack inttrupt, see section 21.6.1 on pg 493 of LPC17XX_UM */
  LPC_TIM0->IR = BIT(0);

  g_timer_count++ ;
}
```

Listing 8.9: Timer0 IRQ Sample Code timer.c

Part III

Questions and Answers

Chapter 9

Frequently Asked Questions

We list frequently asked questions with answers in this chapter.

9.1 Processor Management

This section contains the frequently asked questions and answers of the following topics:

- context switching, and

- preemptive scheduling .

9.1.1 Context Switching

Q1: In tutorial slides, it showed our processes as having 5 states: Ready, Blocked on Resource, Waiting for Message, Executing, and Interrupted. Will these five states be sufficient for our lab, or should we, for example, have a"New" state (as in the context switching sample code on GitHub)?

A1: The NEW state says a process which has never been run before is now ready to run. Tutorial slides are at a higher abstraction level. If you feel adding a new state will help you to manage the processor, then feel free to add more states. The NEW state on GitHub can be considered as a special READY state in tutorial slides.

Q2: If a process completes itself without calling `release_processor()`, should it be taken off the ready queue automatically? Should the next process be dispatched?

A2: The project requires processes never terminate (i.e. they never exit).

Q3: What does the `__set_MSP()` function do, and how is it able to update the program counter/link register to switch to a new process?

A3: This function set the MSP to the value supplied. Each process has its own stack, by switching between stacks, we achieve the context switching purpose.

Q4: How does saving/loading process contexts work in assembly code? Is there any part of the sample code which would help us understand this?

A4: You will need to use assembly code to access C data structures. Section 8.4 in the SE350 lab manual provides some code excerpts on this topic.

9.1.2   Preemptive Scheduling

Q1: What is preemption?

A1: Preemption in this project means when a process whose priority is higher than the priority of the current running process becomes ready, then the current process should be taken out of the processor and let the higher priority ready process to run. In other words, the preemption will guarantee a higher priority ready process will run before a lower priority ready process. It implies that if process P has higher priority than process Q, then at any moment, the system should not allow that P is in READY state and process Q is in RUN state.

Here is a more detailed example: assume we have two processes P and Q. priority(P) > priority(Q). P is blocked on some sort of resource. Q is currently running. Preemption means if an interrupt (software or hardware) happens and it results in a process (say, P) with higher priority than the current running process (say, Q) changing state to READY, then when the interrupt finishes, P will run instead of Q. We say P preempts Q.

Q2: In the project description, when it says a process *may* get preempted, does this imply preemption is not mandatory?

A2: Preemption is mandatory for the entire project and for very single deliverable. However a preemptive API call may not necessarily cause a preemption to happen because it depends on what the values you pass into the API input parameters. For example, `set_process_priority` is a preemptive call. It takes two parameters, a PID and a priority. Calling this function may or may not cause preemption depending on the values of the input parameters of this function call and priorities and states of other processes in the system. Assume we only have two processes P1 and P2. Both

are at MEDIUM priority. P1 is currently in RUN state and P2 is in READY state. Example 1: P1 calls set process priority API to lower down itself to LOW priority. This case, P1 will be preempted by P2. P2 will start to RUN after the API call. Example 2: P1 calls set process priority API to increase itself to HIGH priority. This case, P1 will NOT be pre-empted by P2. P1 will continue in RUN state after the API call.

Q3: Which deliverable requires preemption to be implemented?

A3: Preemption is required for all deliverables (i.e. P1, P2 and P3).

Q4: Since we aren't using timers/interrupts in P1, I was a little unclear what sort of preemption was needed. One method we discussed having each system call call the scheduler, instead of possible returning back to the user code. Can you provide some more details on this though?

A4: Suppose a process calls `set_process_priority` to lower down its own priority, then it might be preempted due to this function call. We do not have hardware interrupts in P1, however we do have software interrupt `SVC`. There are other scenarios a preemption could happen as well in P1.

9.2 Memory Management

Q1: Where do we allocate memory for the blocked resource queue?

A1: You could either allocate the queue in compile time or at run time.

Q2: Do we even need the block resource queue? Can't we use the array of `PCBs` and `PROC_INIT` to get blocked and priority?

A2: You are free to implement the RTX the way you want as long as the final project satisfies the required behavior of the APIs and the processes. You have the freedom to make data structure and algorithm choices, though different choices imply different performance of your RTX. Assume that you do not have a blocked on resource queue, when a certain resource becomes available, you will need to traverse the entire pcb table to find the process that is waiting for the resource and try to unblock it, this probably will be OK if you only have small number of pcbs, but will show performance problem when the number of pcbs grow.

Q3: In our `request_memory_block` function for the kernel, how do we know which process is requesting memory?

A3: Normally the kernel will have a global variable that points to the current running process. The Context Switching sample project names this variable `g_current_process`. So you know it is this process that is requesting the memory.

Q4: If we want to get the end of RAM, instead of hardcoding 0x10008000**, can we use** Image$$RW_IRAM1$$RW$$Limit**? Or some other symbol?**

A4: The end of RAM can only be hard-coded since it is tied to the specific hardware. The `Image$$RW_IRAM1$$RW$$Limit symbol` is put at the end of the Image and it should always be less than `0x10008000`.

Q5: Is each stack element 8 bytes? Or is it 4 bytes?

A5: Each stack element is 4 bytes.

Q6: Why does the stack pointer need to be 8-byte aligned on ARM?

A6: The ARM Architecture Procedure Call Standard (AAPCS) requires stack to be double-word aligned at a public interface. The AAPCS is posted in Learn under manufacture manual section for your reference.

Q7: I'm looking at k_memory.c**. What are these 8-byte alignment shenanigans?**

```
U32 *gp_stack;
/* code to initialize gp_stack skipped */
if ((U32)gp_stack & 0x04) { /* 8 bytes alignment */
    --gp_stack;
}
```

A7: The code basically checks whether the `gp_stack` is aligned by four bytes (i.e. bit AND with binary 100), but not aligned by eight bytes. If yes, then decrement it by four bytes to make it eight bytes aligned. The `gp_stack` is a U32 *, the compiler will give it an address of four bytes aligned. So it is either eight bytes aligned or not eight bytes aligned, but always four bytes aligned. Because it is U32 *, pointer arithmetic `gp_stack` decrements it by four bytes and we get eight byte alignment.

9.3 Interprocess Communication (IPC)

Q1: What is a message envelope?

A1: The project requires a message-based IPC scheme. Messages are carried in shared memory blocks. A process writes a message into a shared memory block, sends a pointer to the memory block to another process and receiving process reads the message from the memory block. The shared memory blocks are called "message envelopes.

Q2: What is the format of the message envelope?

A2: The message envelope contains two data structures. One is for the kernel management purpose. For example it may contain some linked list pointers. The second data structure is exposed to user space and the `struct msgbuf` data structure is defined in Section 2.3. This data structure contains the message type and the actual body of the message.

Q3: Where does the message envelope come from and how does user process use it to send a message?

A3: The user process calls `request_memory_block` primitive to request a memory block to be used as a message envelope and the user will start to fill the `struct msgbuf` data structure defined in Section 2.3. For example, the following user process (wall clock process) code will send a command registration message to process KCD (Keyboard Command Decoder).

```
struct msgbuf *p_msg_env = (struct msgbuf *) request_memory_block();
p_msg_env->mtype = KCD_REG;
p_msg_env->mtext[0] = '%';
p_msg_env->mtext[1] = 'W';
send_message(PID_KCD, (void *)p_msg_env);
```

Q4: The lab manual has the following structure:

```
struct msgbuf {
    int mtype;    /* user defined message type */
    char mtext[1]; /* body of the message */
};
```

The slides have something like this

```
msg t (an envelope) {
    next message
    Sender PID
```

```
    Destination PID
    Message Type
    Message Data
};
```

which one is correct?

A4: The slides give you one possible message envelope data structure as an example to illustrate the idea of a message envelop. Please follow the specification in the lab manual. The manual removed the next message, Sender PID and Destination PID fields from the user process view and you could put them into the kernel pointers part in the message envelope. The manual specification wants to hide more information from the view of a user process.

Q5: The definitions of `receive_message()` are different:

Manual: `void *receive_message(int *sender_id);`

Slides: `msg_t  receive_message();`

which one is correct?

A5: Follow the specification in the manual. Sorry that we missed this piece of outdated information in the slides.

9.4   Processes

9.4.1   The Six User Test Processes

Q1: Do we need to write testing cases to test the six test processes?

A1: The purpose of the user test processes is to implement test cases at user level to test the OS APIs your kernel provides. You do not really test the test processes.

Q2: Does each user test process implement one test case/scenario?

A2: No. Some test scenarios may require more than one test process. One example is to test a process to be blocked on memory when the system is out of memory. This scenario cannot be tested with just one process.

9.4.2 System Processes

Q1: Is system process a kernel process or user process?

A1: A system process can run as a user process or a kernel process depending what the system process is doing and what resources it needs.

Q2: Where should we put the null process (which file)?

A2: Where to put the process in a file is up to you. You have the sole freedom to organize the code in this project. To facilitate the automated third-party user level testing, you should not put the null process in the same file where you define the six user test processes.

Q3: Are system processes such as KCD and CRT scheduled the same as user processes?

A3: Yes.

Q4: Does the sample interrupt-driven UART code include any mechanism for calling our keyboard/CRT routines on interrupt? Or do we need to write that ourselves?

A4: The provided sample UART IRQ code contains both receive and transmit interrupts handling logic. You can borrow the relevant part of the code to write your own CRT, KCD and I-process code.

9.4.3 The I-processes

Q1: Why does an I-process need a pcb?

A1: Because the RTX uses message passing to request interrupt related services and the I-process needs at least a process ID to facilitate this.

Q2: How to invoke an I-process?

A2: The I-process is triggered (or scheduled) by interrupts. When an interrupt fires, it goes to the specific IRQ service routine and this service routine will need to save the context of the current running process (say P) and then makes the code that implement the I-process (most likely a function) to run. Once the I-process finishes, the scheduler may pick another process to run (say Q, P and Q may not necessarily be the same) and the context of Q needs to be restored properly.

There are still context switching involved, but due to the special two state property of an I-process, some students may choose not to do a full fledged context switching as a normal process would do to run the I-process. The bottom line is that an I-process does not get interrupted and the process before the I-process and after the I-process need to have context saved and restored properly.

Q3: Do we need to save the context of the I-process and restore it later?

A3: An i-process only has two states: `RUN` or `WAITING_FOR_INTERRUPT`. It does not get interrupted, so there is no need to save or restore i-process's context. But there is a need to save the context of the process that gets interrupted by the I-process and restore the context of the process that the system decides to run after the I-process finishes.

Q4: To get the interrupt processes to act immediately, I am considering giving them a process priority of -1 so that the scheduler gives interrupt processes the highest priority. This would allow a smooth integration of these processes into the current OS. One of the functions that we must have is `get_process_priority(int processID)` returns -1 on error, and just above that, it says process' can only have priority of 0 - 4. Is this function only expected to work for user processes, thus our solution is viable? Or are we not allowed to use the priority of -1 as specified by the 0-4 range for some reason?

A4: The project specification requires when a user process calls `get_process_priority`, the function returns -1 on error. If you want to have "-1" as a valid process priority for special processes such as i-process, then you have to hide this fact from regular user process. The i-process is *scheduled* by the interrupt, so technically it by-passes the regular process scheduler. However some students in the past did implement the i-process invocation by giving it the highest priority in the system. As long as your implementation follows the specification, we accepted it.

Q5: When an I-process use message passing to send data to the KCD and CRT process, that means they have to request a memory block. What happens if there is no memory in the system? An I-process cannot be memory blocked so should we just disregard that particular interrupt?

A5: In the project description, it mentioned to adapt the behaviour of blocking primitives when they are called by an I-process. You should never let an I-process be blocked.

9.5 Interrupts

Q1: Do we need to use atomic(on) and atomic(off) for the RTX P1 deliverable. If yes can you explain how atomic works?

A1: In sense of this project, the atomic(on) turns on all interrupts and atomic(off) turns off all interrupts. For P1, we assume there is no interrupts, so no need for atomicity for P1. But they are needed in P2 and P3.

Q2: Why do we need to call the scheduler at the end of an interrupt handling process? Wouldn't simply going back to the previously running process work?

A2: The reason is that an interrupt may cause a preemption. For example, if you have a higher priority process blocked on receiving a message that is sent by `delayed_send()` primitive. A timer interrupt may cause the message to be delivered to this higher priority process. Hence when the interrupt finishes, you should not return to the previous process which has lower priority.

Q3: What should the correct behaviour be when a new interrupt happens while handling a previous interrupt?

A3: In this project, nested interrupt is not required. You could just disable all interrupts when you are inside the interrupt handler.

Q4: How to enable and disable all interrupts?

A4: This can be done either in the assembly code or the C code. Please see table 8.3.

Q5: Is the TX interrupt activated by setting the TX FIFO rest bit specified in Table 278 on p305 of LPC17xx user's manual[4]?

A5: It is the THRE Interrupt Enable bit in Table 275 on page 302.

9.6 Keil IDE

Q1: Is there a way we can have two configurations, one for debug and one for release? This way we can have a DEBUG macro be defined only in the debug configuration.I noticed you use `DEBUG_0`. Is this some special macro?

A1: The conditional compilation symbol `DEBUG_0` is defined in target options under C/C++ tab. This is a symbol we define for debugging purpose which makes the code to print out more debugging messages. You could define your own symbol to control compilation in a similar way.

Q2: Is it possible to trigger UART interrupts via the keyboard while debugging with the simulator?

A2: Yes, just type in the `UART#1` window in simulator (assuming you are programming UART0). If you are looking at the sample UART IRQ code on GitHub, you need to work on the SIM target in order to see IRQs under simulator. Not the RAM target.

Appendix A

MDK-ARM Installation

There is only a windows port for the Keil MDK-ARM for now. The MDK-ARM V4.60.0.0 direct download link is inside the Learn (`http://learn.uwaterloo.ca`) [1].

During the process of the installation of the MDK-ARM, you will be asked to add example code. Choose Keil(NXP) MCB1xxx Boards example projects (see Figure A.1).
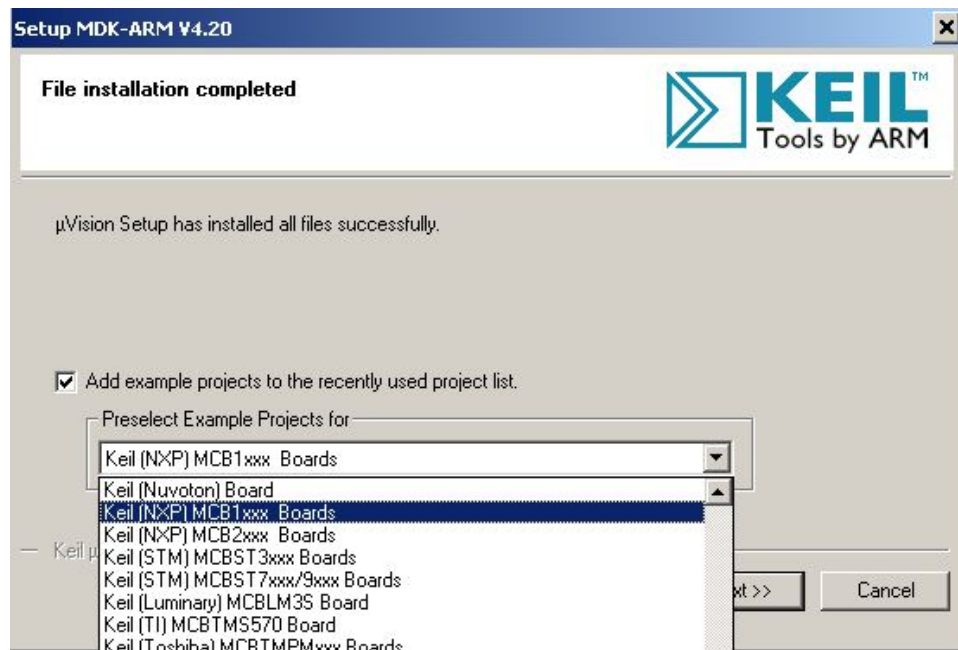

Figure A.1: MDK-ARM Installation Steps: Choose Example Projects

---

[1]The latest version of MDK-ARM is at Keil website http://www.keil.com/download/product/. However the lab manual is written for V4.60.0.0. We haven't tested the latest version.
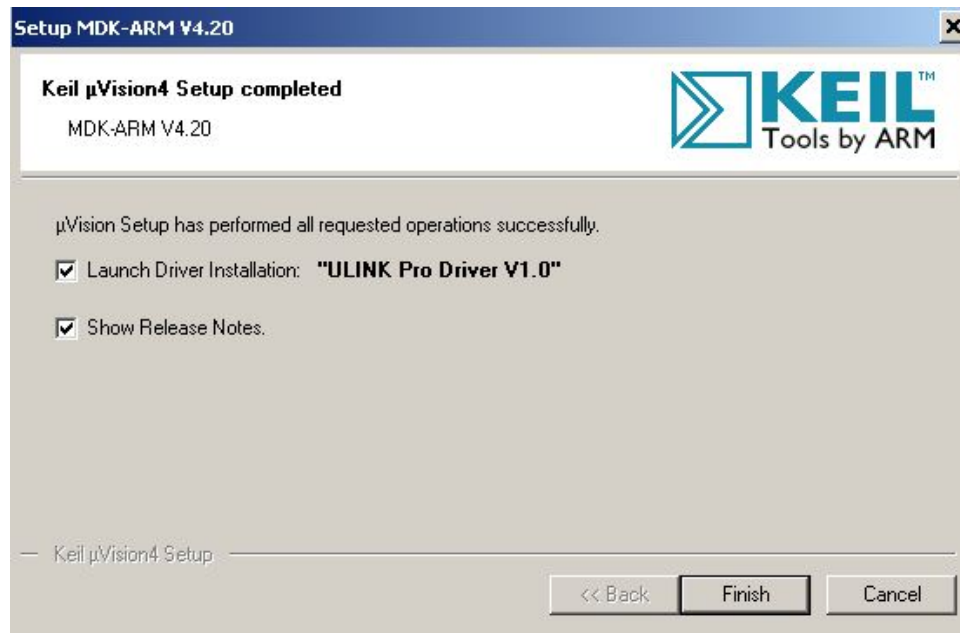
Figure A.2: MDK-ARM Installation Steps: Finish

At the last step of MDK-ARM installation, be sure that the launch the "ULINK Pro Driver

Figure A.3: MDK-ARM Installation Steps: ULINK Pro Driver

V1.0" driver installation check box is checked (see Figure A.2. Once you click "Finish" button, the ULINK Pro Driver installation starts. Click "Install" button to install the driver (see Figure A.3).

Appendix B

Forms

Lab administration related forms are given in this appendix.

SE350 Request to Leave a Project Group Form

Name:	
Quest ID:	
Student ID:	
Lab Assignment ID	
Group ID:	
Names of Other Group Members:	

Provide the reason for leaving the project group here:

Signature _____   Date _____

Bibliography

[1] MCB1700 User's Guide. http://www.keil.com/support/man/docs/mcb1700. 24

[2] MDK Primer. http://www.keil.com/support/man/docs/gsac. 49, 50, 51

[3] Realview compilation tools version 4.0: Compiler reference guide, 2007-2010.

[4] LPC17xx User Manual, Rev2.0, 2010. 23, 26, 27, 30, 42, 46, 52, 55, 69, 82

[5] J. Yiu. *The Definitive Guide to the ARM Cortex-M3*. Newnes, 2009. 26, 28, 32, 33, 49